Interval search experiments

Built 2026-10-17 02:05 PT. Experiment design by Roland Gao; runs and analysis by Claude. Code: research/cooldown_search_probe.py (runs), research/cooldown_study.py (analysis), research/make_cooldown_study.py (this PDF). Logs: logs/research_cooldown_search_probe_*, logs/research_hparam_importance_*. Batch size 2000 throughout.

Experiment 1. Cooldown study: which 40-step lookahead ranks interval candidates like the final result?

Question

An interval search picks each interval's LR by training the interval with a candidate LR, then a short lookahead (the cooldown), and scoring the result. Which cooldown ranks the candidates the way the final accuracy does? Compared, all with a 40-step lookahead (one interval at batch size 2000): no lookahead; linear decay to 0 (conv only, or all four groups); one searched constant cooldown, either for conv only with the other groups on their ground-truth schedules, or shared by all four groups; and four searched constant cooldowns (one per group).

Setup

Item	Value
Model / data	CIFAR-10 airbench, batch size 2000, 200 steps, baseline_bs2000 (conv Muon2 0.22 -> 0; head, norm_bias, whiten_bias SGD on their baseline schedules)
Intervals	4 intervals of 40 steps: steps 0-40, 40-80, 80-120, 120-160 (the last interval has no room for a 40-step lookahead)
Candidates	7 constant conv LRs for the interval: the 0.6^k grid point nearest the baseline's mean conv LR in that interval, and 3 grid steps either side; other groups on the baseline
Ground truth	Final test TTA when training continues on the baseline schedule after the interval
No lookahead	Test TTA right after the interval
Linear to 0	Conv only: conv decays linearly from the candidate to 0 over 40 steps, the other groups continue their schedules. All groups: every group decays linearly to 0 from its current value
1 searched cooldown: conv only	Conv constant at m x candidate for 40 steps; the other groups follow their baseline (ground-truth) schedules; m chosen from {1, 0.6, 0.36, 0.22, 0.13, 0.078, 0.047, 0} by the best score on that seed
1 searched cooldown: shared	Constant LRs for 40 steps: conv at m x candidate, each other group at m x its mean LR over the interval; one m for all groups, chosen as above
4 searched cooldowns	As above but one m per group; coordinate search from the best shared m, one group at a time over the same 8 values, until no group changes (at most 3 passes)
Seed	1 training seed (0); every method uses the same prefix and data order
Score	Test-set TTA accuracy (%)

Results

Picked LR = the candidate with the highest score. Regret = ground truth at its best candidate minus ground truth at the picked one (pp). One seed: single-run noise is about 0.13 pp, so differences of that size between candidates are not resolved.

Picked LR (regret, pp)

Method	Interval 1	Interval 2	Interval 3	Interval 4
Ground truth (final TTA)	0.078 (0.00)	0.13 (0.00)	0.078 (0.00)	0.078 (0.00)
No lookahead	0.078 (0.00)	0.028 (0.34)	0.028 (0.01)	0.028 (0.29)
Linear to 0: conv only	0.13 (0.05)	0.078 (0.26)	0.078 (0.00)	0.028 (0.29)
Linear to 0: all four groups	0.13 (0.05)	0.078 (0.26)	0.047 (0.21)	0.028 (0.29)
1 searched cooldown: conv only, others on ground truth	0.13 (0.05)	0.078 (0.26)	0.047 (0.21)	0.047 (0.21)
1 searched cooldown: shared by all four groups	0.13 (0.05)	0.078 (0.26)	0.047 (0.21)	0.047 (0.21)
4 searched cooldowns (one per group)	0.13 (0.05)	0.078 (0.26)	0.047 (0.21)	0.047 (0.21)

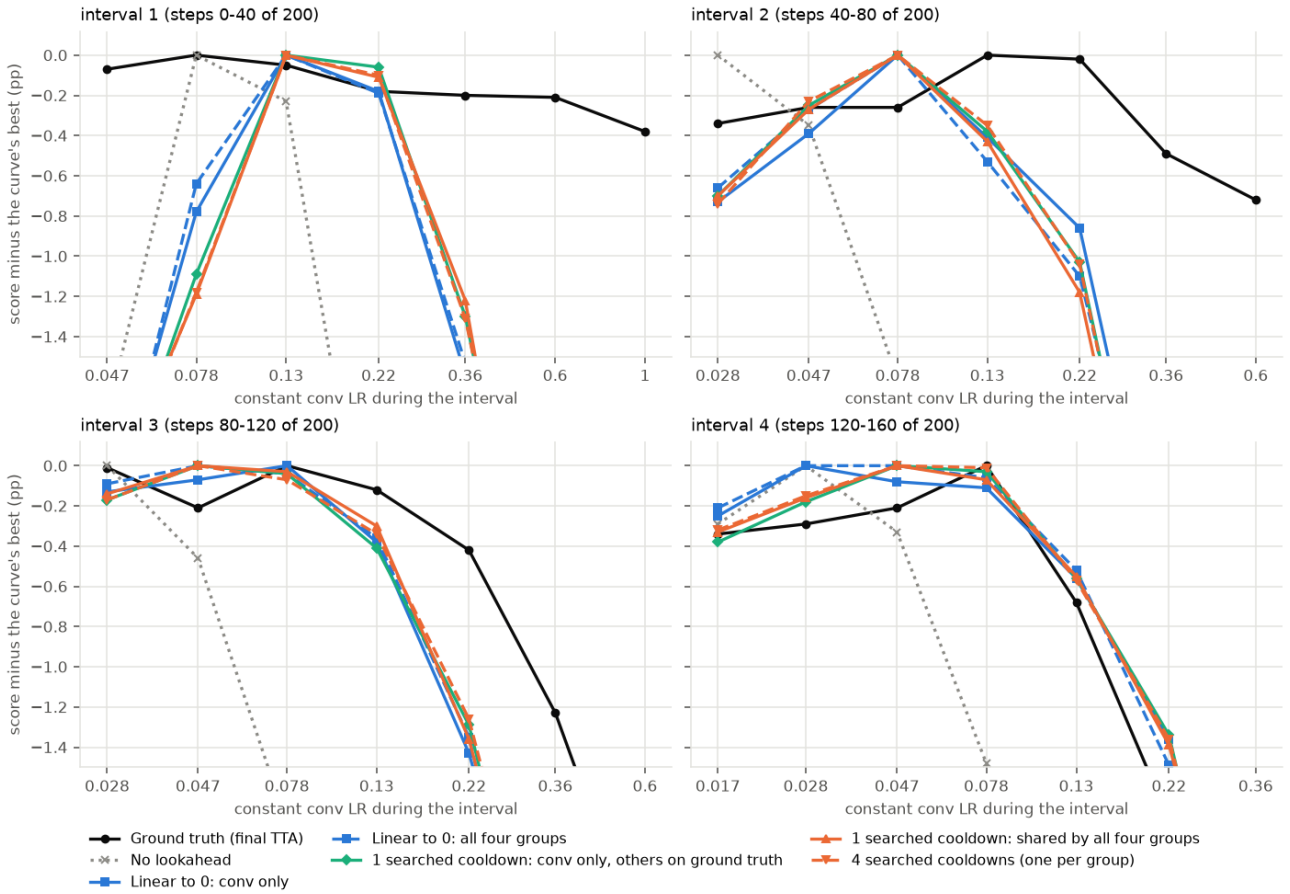


Figure 1. Each method's score per candidate, relative to its own best (the point at 0 is its pick).

Interval 1: score per candidate conv LR (bold: each row's best)

Method	0.047	0.078	0.13	0.22	0.36	0.6	1
Ground truth (final TTA)	93.99	**94.06**	94.01	93.88	93.86	93.85	93.68
No lookahead	81.57	**83.35**	83.12	80.45	68.80	56.07	35.18
Linear to 0: conv only	89.04	90.59	**91.37**	91.19	89.78	87.38	84.39
Linear to 0: all four groups	88.85	90.63	**91.27**	91.08	89.74	87.41	84.10
1 searched cooldown: conv only, others on ground truth	88.58	89.83	**90.92**	90.86	89.62	87.45	85.95
1 searched cooldown: shared by all four groups	88.65	89.70	**90.89**	90.78	89.67	87.16	85.13
4 searched cooldowns (one per group)	88.79	89.89	**91.07**	90.97	89.76	87.65	86.13

Candidate	Conv only: chosen m (seeds)	Shared: chosen m (seeds)	4 searched: chosen (conv, head, norm_bias, whiten_bias) (seeds)
0.047	0.6 x1	0.6 x1	(1, 1, 0.6, 0.36) x1
0.078	0.36 x1	0.36 x1	(0.36, 1, 1, 0.22) x1
0.13	0.22 x1	0.22 x1	(0.22, 1, 0.22, 0.22) x1
0.22	0.22 x1	0.22 x1	(0.22, 0.36, 0.22, 1) x1
0.36	0.13 x1	0.13 x1	(0.13, 0.047, 0.13, 0.13) x1
0.6	0.078 x1	0.078 x1	(0.078, 1, 0.36, 0.078) x1
1	0.078 x1	0.078 x1	(0.078, 0.6, 0.078, 0.6) x1

Interval 2: score per candidate conv LR (bold: each row's best)

Method	0.028	0.047	0.078	0.13	0.22	0.36	0.6
Ground truth (final TTA)	93.76	93.84	93.84	**94.10**	94.08	93.61	93.38
No lookahead	**91.12**	90.77	89.51	87.43	82.45	71.27	57.08
Linear to 0: conv only	92.25	92.59	**92.98**	92.57	92.12	90.18	87.94
Linear to 0: all four groups	92.27	92.66	**92.93**	92.40	91.83	90.11	87.78
1 searched cooldown: conv only, others on ground truth	92.16	92.61	**92.86**	92.48	91.83	89.82	87.93
1 searched cooldown: shared by all four groups	92.18	92.61	**92.88**	92.45	91.70	89.73	87.84
4 searched cooldowns (one per group)	92.20	92.71	**92.94**	92.59	91.90	89.87	87.98

Candidate	Conv only: chosen m (seeds)	Shared: chosen m (seeds)	4 searched: chosen (conv, head, norm_bias, whiten_bias) (seeds)
0.028	0.6 x1	0.6 x1	(0.6, 0.6, 1, 0.6) x1
0.047	0.36 x1	0.36 x1	(0.36, 1, 0.047, 0.36) x1
0.078	0.22 x1	0.22 x1	(0.22, 0.13, 0.22, 0.22) x1
0.13	0.22 x1	0.22 x1	(0.22, 1, 0.36, 0) x1
0.22	0.13 x1	0.13 x1	(0.13, 1, 1, 0.078) x1
0.36	0.13 x1	0.13 x1	(0.13, 0.6, 0, 0.047) x1
0.6	0.078 x1	0.078 x1	(0.078, 0.047, 0.13, 0.078) x1

Interval 3: score per candidate conv LR (bold: each row's best)

Method	0.028	0.047	0.078	0.13	0.22	0.36	0.6
Ground truth (final TTA)	93.99	93.79	**94.00**	93.88	93.58	92.77	91.54
No lookahead	**92.59**	92.13	90.81	87.68	82.32	71.78	56.20
Linear to 0: conv only	93.26	93.32	**93.39**	93.01	91.96	90.45	87.63
Linear to 0: all four groups	93.28	**93.37**	93.33	93.01	92.02	90.23	87.59
1 searched cooldown: conv only, others on ground truth	93.16	**93.33**	93.29	92.92	92.04	90.31	88.61
1 searched cooldown: shared by all four groups	93.13	**93.27**	93.24	92.97	91.91	90.29	88.64
4 searched cooldowns (one per group)	93.18	**93.35**	93.28	93.01	92.09	90.44	88.76

Candidate	Conv only: chosen m (seeds)	Shared: chosen m (seeds)	4 searched: chosen (conv, head, norm_bias, whiten_bias) (seeds)
0.028	0.36 x1	0.22 x1	(0.22, 0.13, 1, 0.22) x1
0.047	0.22 x1	0.36 x1	(0.22, 1, 1, 0.36) x1
0.078	0.22 x1	0.22 x1	(0.22, 1, 0.36, 0.22) x1
0.13	0.22 x1	0.22 x1	(0.22, 1, 0.22, 0.22) x1
0.22	0.22 x1	0.13 x1	(0.22, 0.36, 1, 0.13) x1
0.36	0.13 x1	0.13 x1	(0.13, 0.078, 0.6, 0.13) x1
0.6	0.078 x1	0.13 x1	(0.078, 0.36, 0.13, 0.13) x1

Interval 4: score per candidate conv LR (bold: each row's best)

Method	0.017	0.028	0.047	0.078	0.13	0.22	0.36
Ground truth (final TTA)	93.65	93.70	93.78	**93.99**	93.31	92.23	89.58
No lookahead	92.89	**93.18**	92.85	91.70	88.84	83.30	72.23
Linear to 0: conv only	93.55	**93.80**	93.72	93.69	93.24	92.44	90.63
Linear to 0: all four groups	93.55	**93.76**	**93.76**	93.70	93.24	92.27	90.57
1 searched cooldown: conv only, others on ground truth	93.51	93.71	**93.89**	93.86	93.33	92.55	90.70
1 searched cooldown: shared by all four groups	93.56	93.73	**93.89**	93.82	93.34	92.50	90.89
4 searched cooldowns (one per group)	93.62	93.79	**93.94**	93.93	93.37	92.58	90.93

Candidate	Conv only: chosen m (seeds)	Shared: chosen m (seeds)	4 searched: chosen (conv, head, norm_bias, whiten_bias) (seeds)
0.017	0.6 x1	0.6 x1	(0.6, 1, 1, 0.6) x1
0.028	0.6 x1	0.6 x1	(0.6, 1, 0.36, 0.6) x1
0.047	0.36 x1	0.36 x1	(0.36, 0.22, 1, 0.36) x1
0.078	0.22 x1	0.22 x1	(0.22, 0.047, 0.22, 0.22) x1
0.13	0.13 x1	0.13 x1	(0.13, 0.22, 0.13, 0.13) x1
0.22	0.13 x1	0.13 x1	(0.13, 1, 0.078, 0.13) x1
0.36	0.13 x1	0.13 x1	(0.13, 0.078, 1, 0.13) x1

Observations

One seed: single-run noise is about 0.13 pp, so ground-truth differences of that size are not resolved. Interval 1's ground truth is flat (93.85-94.06 from 0.047 to 0.6) and interval 3's is non-monotone (0.028: 93.99, 0.047: 93.79, 0.078: 94.00), so regrets of about 0.2 pp or less at those intervals are within noise.

All five 40-step lookahead variants pick the same LR at intervals 1 and 2, and the three searched-constant variants (conv only, shared, four per group) pick the same LR at every interval. How the other three groups cool down (continuing their ground-truth schedules, one shared multiplier, or their own multipliers) never changed the pick. Four searched multipliers score slightly higher in absolute terms than one but rank the candidates the same way.

The 40-step lookahead picks a smaller LR than the ground truth where the ground truth is clearly peaked: interval 2 (all variants pick 0.078 where 0.13 is best, 0.26 pp regret) and interval 4 (0.028-0.047 where 0.078 is best, 0.21-0.29 pp). No lookahead is worse at intervals 2 and 4 (0.34 / 0.29 pp) and always picks the smallest or second-smallest candidate after interval 1.

Linear decay and searched constants differ only at intervals 3 and 4, in opposite directions (linear conv-only is better at interval 3, searched constants at interval 4), both by about 0.2-0.3 pp: not resolved with one seed.

Experiment 2. Hparam importance from scratch

Question

How important is each of the four LRs when they are searched jointly from scratch, and does the answer depend on the main interval length?

Setup

Item	Value
Model / data	baseline_bs2000 except the four trained LRs (conv, head, norm_bias, whiten_bias), which are constant for the first L training steps from initialization
Main interval L	1, 2, 5, 10, 40 steps; the score is taken right after the L steps (no lookahead)
Joint search	From scratch: every LR starts at 1.0. Scan each LR over x10^k, k = -3..3 (others fixed), keep the best, snap to the 0.6^k grid, then coordinate-climb on the grid (one step up / down, move while the score improves) until no LR moves: the center
Sweep	From the center, each LR alone over 5 grid points on each side (others at the center). Offset +1 = one grid step larger (x1/0.6), -1 = one step smaller (x0.6)
Seed / score	1 seed (0); test-set TTA accuracy (%); the test cross-entropy loss is logged too

Joint-search centers

L (steps)	Center: conv, head, norm_bias, whiten_bias	Center TTA	Evaluations
1	0.13, 60, 0.00078, 60	25.43	73
2	0.22, 1, 1.7, 1300	36.89	71
5	0.078, 0.006, 2100, 1300	52.00	76
10	0.078, 0.017, 2100, 13	59.47	70
40	0.078, 2100, 280, 2100	84.60	73

Importance: TTA drop from the center (pp) at 1 step / 2 steps (mean of both sides) / worst within 5 steps

LR	L = 1	L = 2	L = 5	L = 10	L = 40
conv	0.58 / 2.01 / 16.42	1.00 / 2.90 / 19.69	3.20 / 7.96 / 32.96	2.02 / 4.11 / 29.37	1.95 / 5.26 / 49.82
head	0.03 / 0.14 / 0.29	0.02 / 0.02 / 0.08	0.10 / 0.12 / 0.18	0.32 / 0.49 / 0.67	0.54 / 1.65 / 62.56
norm_bias	0.00 / 0.00 / 0.00	0.05 / 0.02 / 0.08	0.51 / 1.46 / 28.09	1.12 / 2.77 / 28.84	0.43 / 0.84 / 6.98
whiten_bias	0.01 / 0.03 / 0.09	0.15 / 0.85 / 4.48	0.56 / 1.02 / 3.45	0.51 / 0.58 / 2.00	0.46 / 0.96 / 2.44

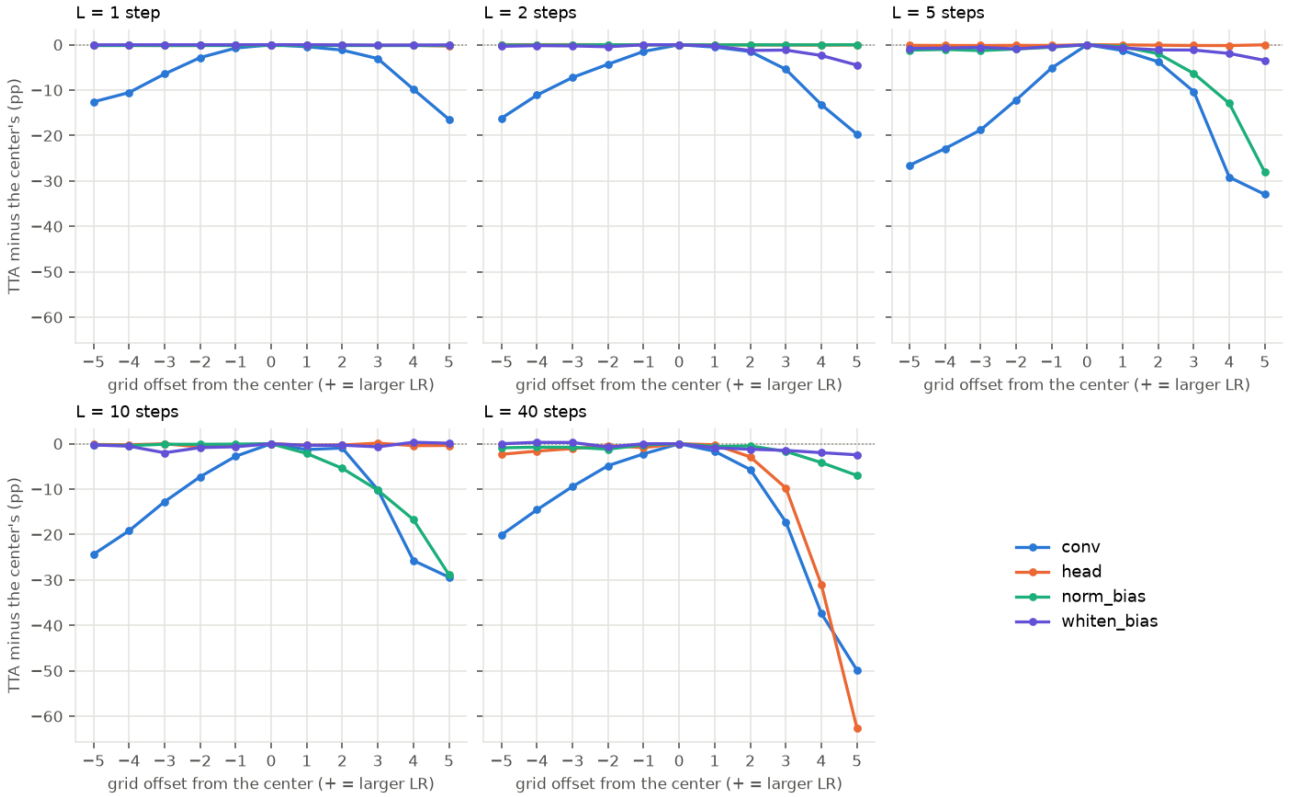


Figure 2. One-at-a-time sweeps around each L's joint-search center.

L = 1: TTA minus the center's (pp) at each grid offset (center LR in brackets; center TTA 25.43)

LR	-5	-4	-3	-2	-1	+0	+1	+2	+3	+4	+5
conv (0.13)	-12.58	-10.53	-6.44	-2.84	-0.73	+0.00	-0.43	-1.18	-3.06	-9.80	-16.42
head (60)	-0.15	-0.14	-0.17	-0.16	-0.03	+0.00	-0.03	-0.11	-0.16	-0.12	-0.29
norm_bias (0.00078)	+0.00	+0.00	+0.00	+0.00	+0.00	+0.00	+0.00	+0.00	+0.00	+0.00	+0.00
whiten_bias (60)	-0.03	-0.02	-0.01	-0.01	-0.01	+0.00	-0.01	-0.05	-0.09	-0.08	-0.04

L = 2: TTA minus the center's (pp) at each grid offset (center LR in brackets; center TTA 36.89)

LR	-5	-4	-3	-2	-1	+0	+1	+2	+3	+4	+5
conv (0.22)	-16.19	-11.09	-7.21	-4.29	-1.48	+0.00	-0.52	-1.50	-5.41	-13.19	-19.69
head (1)	+0.01	+0.01	+0.02	+0.00	-0.01	+0.00	-0.03	-0.03	-0.05	-0.08	-0.04
norm_bias (1.7)	-0.07	-0.08	-0.05	-0.01	-0.04	+0.00	-0.05	-0.03	-0.03	-0.02	-0.01
whiten_bias (1300)	-0.34	-0.20	-0.26	-0.44	-0.01	+0.00	-0.29	-1.26	-1.18	-2.37	-4.48

L = 5: TTA minus the center's (pp) at each grid offset (center LR in brackets; center TTA 52.00)

LR	-5	-4	-3	-2	-1	+0	+1	+2	+3	+4	+5
conv (0.078)	-26.55	-22.87	-18.78	-12.17	-5.13	+0.00	-1.28	-3.75	-10.31	-29.17	-32.96
head (0.006)	-0.13	-0.17	-0.17	-0.14	-0.16	+0.00	-0.04	-0.10	-0.17	-0.18	-0.02
norm_bias (2100)	-1.19	-1.04	-1.29	-0.94	-0.54	+0.00	-0.47	-1.97	-6.33	-12.92	-28.09
whiten_bias (1300)	-0.81	-0.75	-0.58	-0.91	-0.41	+0.00	-0.71	-1.13	-1.17	-1.92	-3.45

L = 10: TTA minus the center's (pp) at each grid offset (center LR in brackets; center TTA 59.47)

LR	-5	-4	-3	-2	-1	+0	+1	+2	+3	+4	+5
conv (0.078)	-24.35	-19.17	-12.77	-7.28	-2.76	+0.00	-1.28	-0.95	-10.13	-25.72	-29.37
head (0.017)	-0.17	-0.23	-0.02	-0.67	-0.31	+0.00	-0.34	-0.30	+0.13	-0.42	-0.38
norm_bias (2100)	-0.29	-0.39	-0.12	-0.13	-0.09	+0.00	-2.15	-5.42	-10.23	-16.70	-28.84
whiten_bias (13)	-0.26	-0.53	-2.00	-0.85	-0.67	+0.00	-0.36	-0.32	-0.68	+0.32	+0.11

L = 40: TTA minus the center's (pp) at each grid offset (center LR in brackets; center TTA 84.60)

LR	-5	-4	-3	-2	-1	+0	+1	+2	+3	+4	+5
conv (0.078)	-20.08	-14.54	-9.42	-4.83	-2.24	+0.00	-1.67	-5.69	-17.22	-37.36	-49.82
head (2100)	-2.33	-1.63	-1.07	-0.40	-0.85	+0.00	-0.23	-2.90	-9.74	-31.04	-62.56
norm_bias (280)	-0.90	-0.75	-0.77	-1.20	-0.24	+0.00	-0.62	-0.48	-1.69	-4.15	-6.98
whiten_bias (2100)	+0.00	+0.30	+0.27	-0.72	-0.01	+0.00	-0.91	-1.21	-1.47	-1.97	-2.44

Experiment 3. Representation quality of schedule families

Question

How close to the baseline can each family of conv LR schedules get, under no transform and under the ramp, reverse ramp and wave transforms, when searched by the same global search?

Setup

Item	Value
Model / data	baseline_bs2000; only the conv LR schedule is searched, the other groups stay on the baseline. The transform multiplies the conv LR the search specifies
Conditions	No transform; ramp x 4^(2x-1) (x0.25 -> x4); reverse ramp x 4^(1-2x) (x4 -> x0.25); wave x 2^sin(2 pi x) (x1 -> x2 -> x0.5 -> x1); x = fraction of training
Families	5 constant intervals (5 values); linear decay to 0 (initial LR); power decay lr x (1-t)^p (lr, p; p = 1 is linear); 2 knots at t = 0, 0.5; 3 knots at t = 0, 1/3, 2/3; 5 knots at t = 0, 0.2, 0.4, 0.6, 0.8 (knot values; linear between knots, 0 at the end)
Search	Generalized global_neighbour (global_family): coordinate search over the family's parameters on multiplicative grids (LR values x0.6, p x0.8), one grid step at a time, moving while the score improves; every candidate is a full training run on seed 0 from the same initial state
Start	The oracle spec (baseline conv schedule divided by the transform) projected onto the family: interval means or knot values snapped to the grid; linear / power start at 0.22 with p = 1
Evaluation	Start and best schedules retrained on held-out seeds 100-107 (mean test TTA; not part of the search)

References (held-out TTA, seeds 100-107)

	No transform	Ramp	Reverse ramp	Wave
Baseline linear decay, no transform (the target)	93.92	93.92	93.92	93.92
Baseline spec kept under the transform (do-nothing)	93.92	93.42	93.57	93.72

Held-out TTA of the best schedule (start = the projected oracle before search; search cost)

Family	No transform	Ramp	Reverse ramp	Wave
5 constant intervals	**93.62** (start 93.52; 18 runs, 182 s)	**93.65** (start 93.32; 28 runs, 285 s)	**93.86** (start 93.81; 28 runs, 279 s)	**93.82** (start 93.11; 22 runs, 222 s)
Linear decay (1 param)	**93.86** (start 93.92; 3 runs, 31 s)	**93.42** (start 93.42; 3 runs, 31 s)	**93.59** (start 93.57; 4 runs, 41 s)	**93.91** (start 93.72; 3 runs, 31 s)
Power decay (lr, p)	**93.86** (start 93.92; 8 runs, 81 s)	**93.95** (start 93.42; 14 runs, 141 s)	**93.59** (start 93.57; 6 runs, 61 s)	**93.91** (start 93.72; 5 runs, 51 s)
2 knots (t = 0, 0.5)	**93.86** (start 93.86; 5 runs, 51 s)	**93.77** (start 93.14; 9 runs, 91 s)	**93.88** (start 93.85; 8 runs, 83 s)	**93.83** (start 93.72; 12 runs, 121 s)
3 knots (t = 0, 1/3, 2/3)	**93.86** (start 93.94; 12 runs, 121 s)	**93.68** (start 93.74; 10 runs, 101 s)	**93.78** (start 93.90; 13 runs, 131 s)	**94.06** (start 93.90; 7 runs, 73 s)
5 knots (t = 0, 0.2, ..., 0.8)	**93.94** (start 93.97; 13 runs, 131 s)	**93.96** (start 93.91; 17 runs, 170 s)	**93.91** (start 93.91; 11 runs, 111 s)	**93.94** (start 93.97; 11 runs, 111 s)

Best parameter values (LR values in order; power: lr, p)

Family	No transform	Ramp	Reverse ramp	Wave
5 constant intervals	0.13, 0.22, 0.078, 0.047, 0.028	0.36, 0.13, 0.078, 0.028, 0.0036	0.028, 0.078, 0.13, 0.078, 0.047	0.078, 0.078, 0.13, 0.078, 0.017
Linear decay (1 param)	0.13	0.22	0.078	0.13
Power decay (lr, p)	0.13, 0.64	0.36, 2	0.078, 1	0.13, 1
2 knots (t = 0, 0.5)	0.22, 0.13	0.6, 0.078	0.078, 0.078	0.13, 0.078
3 knots (t = 0, 1/3, 2/3)	0.13, 0.22, 0.047	0.36, 0.13, 0.047	0.078, 0.078, 0.047	0.13, 0.078, 0.13
5 knots (t = 0, 0.2, ..., 0.8)	0.13, 0.13, 0.13, 0.078, 0.047	1, 0.36, 0.22, 0.047, 0.017	0.047, 0.078, 0.078, 0.13, 0.13	0.13, 0.078, 0.078, 0.13, 0.078

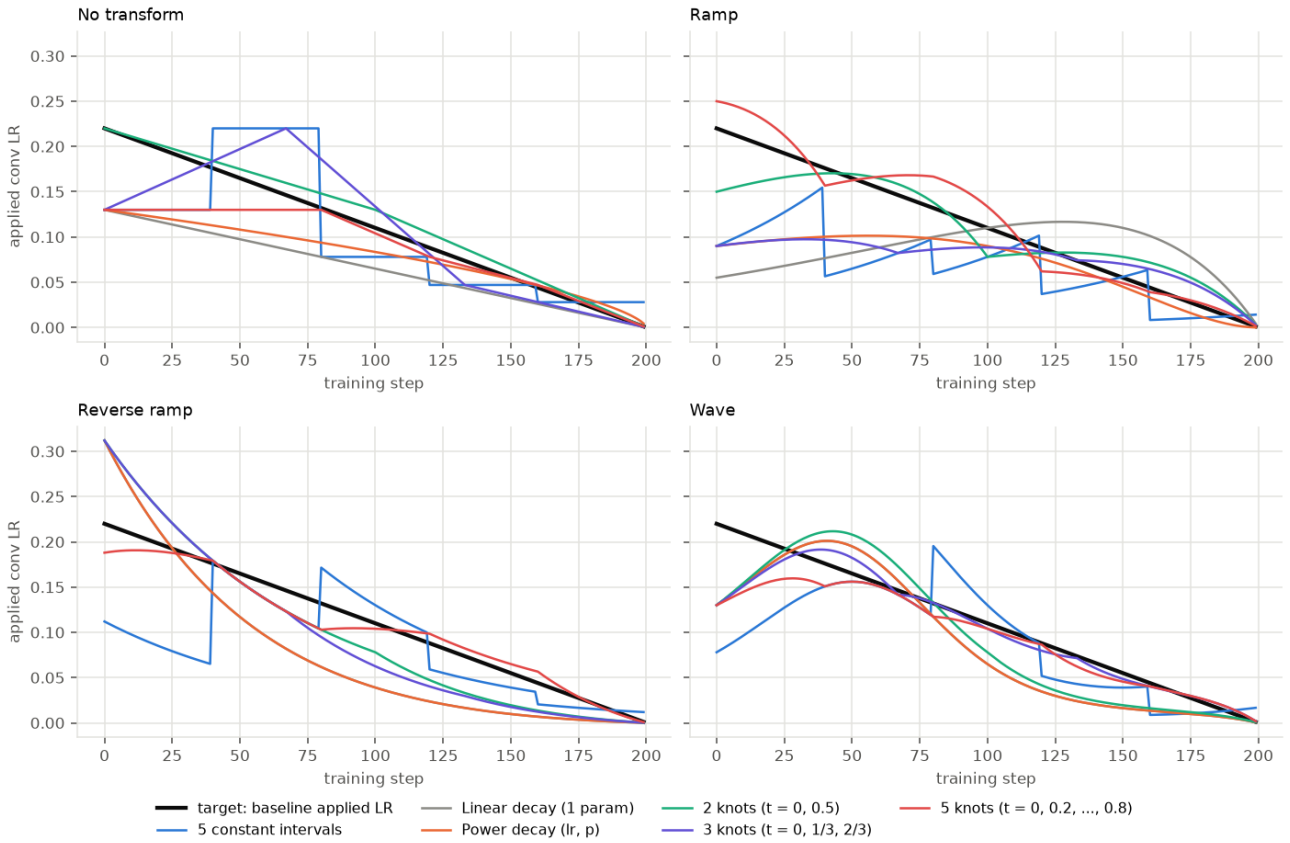


Figure 3. Applied conv LR (spec x transform) of each family's best schedule; the black line is the baseline's applied LR, which a perfect search would reproduce.

Observations

Held-out means over 8 seeds have a standard error of about 0.045 pp, so differences below about 0.1 pp are not resolved. Target: the baseline's 93.92 on the same seeds.

5 knots is the only family that reached the target in all four conditions in these searches (93.91-93.96), and its start, the oracle projected onto the family, is already there (93.91-93.97): 5 knots on the grid represent all four target schedules.

5 constant intervals of conv reach only 93.62 here without a transform (start 93.52), but that is this search's local optimum, not the family's limit. Earlier 5-interval searches at batch size 2000 (same held-out seeds) did better: conv only 93.77 (interval search; conv 0.13, 0.13, 0.078, 0.047, 0.01) and 93.72 (global_interval), and with all four LRs as 5 intervals 93.92 (interval search), matching the baseline, and 93.85 (global_interval). This search kept the last interval at 0.028 where those went to 0.01-0.017: a single-seed coordinate search can stop one grid step short at the end. Under the transforms it reaches 93.65 / 93.86 / 93.82 (ramp / reverse ramp / wave; projected starts 93.32 / 93.81 / 93.11).

Linear decay matches the target without a transform (93.86-93.92) and under wave (93.91), but no initial LR compensates a ramp (93.42) or a reverse ramp (93.59): the shape is wrong, not the scale. Power decay fixes the ramp (93.95 with p = 2, lr 0.36: decaying faster) but not the reverse ramp (93.59): (1-t)^p cannot rise early, which the reverse ramp needs.

2 and 3 knots land in between (93.68-94.06), with no consistent ordering across conditions.

Single-seed search sometimes leaves a good start for a worse held-out schedule (no transform: linear 93.92 -> 93.86, 3 knots 93.94 -> 93.86, 5 constant intervals moved its first value from 0.22 to 0.13): the search optimizes seed 0's test score, the same winner's curse as before.

Experiment 4. Hparam range from scratch (Experiment 2 redone)

Question

Experiment 2's +-5 sweeps depend on where the center landed. How far can each LR move from the center, on each side, before the score drops by 10%? A range that ends on both sides marks an LR with an optimum; a side that never drops marks an LR the score is insensitive to in that direction.

Setup

Item	Value
Model / score	As in Experiment 2: baseline_bs2000 except the four LRs, constant for the first L steps from initialization; test TTA right after the L steps; 1 seed (0)
Center	Experiment 2's joint-search center for each L (re-evaluated; it reproduces Experiment 2's TTA)
Probe	One LR at a time (others at the center), one 0.6^k grid step at a time on each side (+1 = x1/0.6, larger), until TTA <= 0.9 x the center's TTA (a 10% relative drop) or 20 steps. A diverged run counts as a drop

Stop offsets (smaller side / larger side) and the LR range within 10% of the center

LR	L = 1	L = 2	L = 5	L = 10	L = 40
conv	-2 / +3 0.078 - 0.36	-2 / +3 0.13 - 0.6	-2 / +3 0.047 - 0.22	-2 / +3 0.047 - 0.22	-3 / +3 0.028 - 0.22
head	<-20 / +7 0.0022 - 1300	<-20 / +14 3.7e-05 - 770	<-20 / >20 2.2e-07 - 170	<-20 / >20 6.1e-07 - 460	<-20 / +3 0.078 - 5900
norm_bias	<-20 / >20 2.9e-08 - 21	<-20 / +17 6.1e-05 - 5900	<-20 / +3 0.078 - 5900	<-20 / +3 0.078 - 5900	<-20 / +6 0.01 - 3500
whiten_bias	<-20 / +10 0.0022 - 5900	<-20 / +5 0.047 - 9800	<-20 / +6 0.047 - 16000	<-20 / +17 0.00047 - 46000	<-20 / +11 0.078 - 350000

<-20 / >20: no 10% drop within 20 steps on that side (the range shown then ends at the last probed point). Center TTA: L = 1: 25.43, L = 2: 36.89, L = 5: 52.00, L = 10: 59.47, L = 40: 84.60.

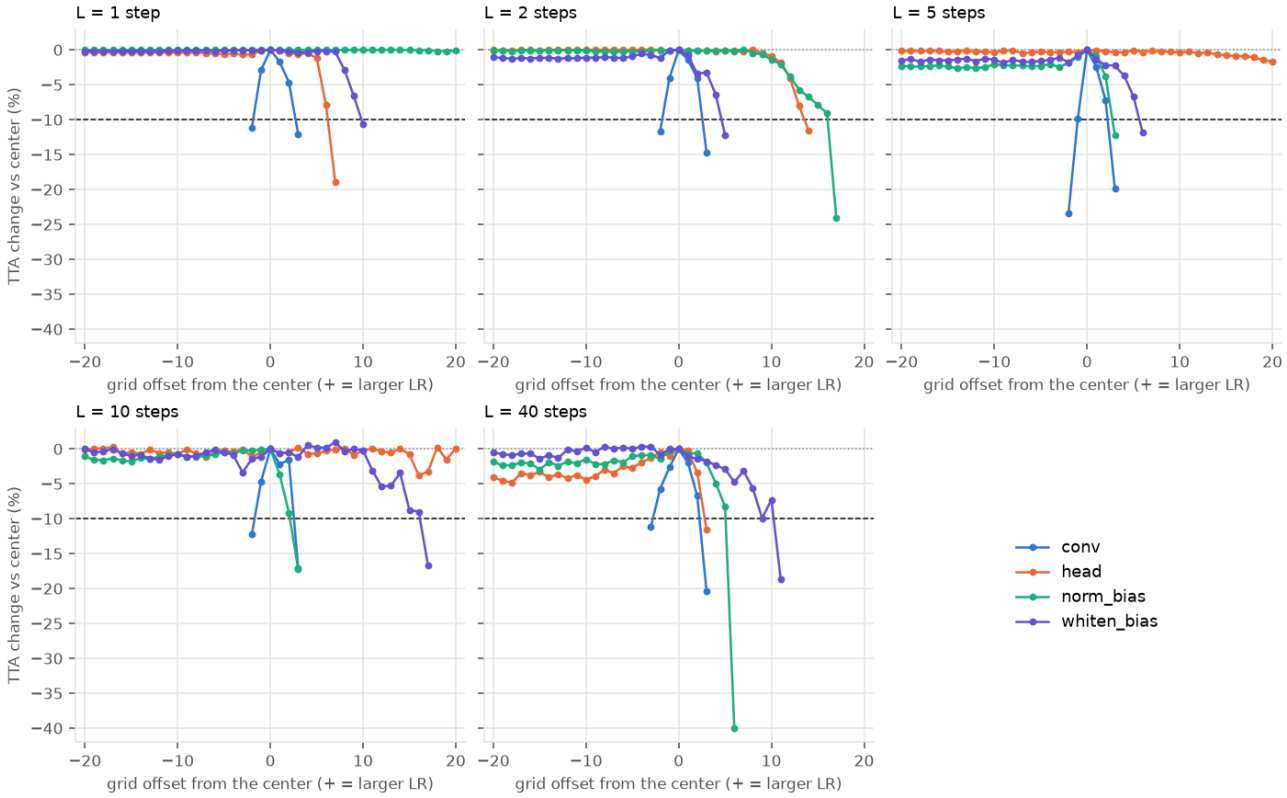


Figure 4. Relative TTA change along each probe; the dashed line is the 10% stop. Points below -40% are drawn at -40%.

Observations

One seed: single-run noise is about 0.13 pp at the full 200 steps; the stop criterion (a 10% relative drop: 2.5 pp at L = 1, 8.5 pp at L = 40) is far above it, so the stop offsets are not noise.

Conv is the only LR with an optimum: its range ends on both sides at every L, at -2 or -3 grid steps (smaller) and +3 (larger). The LR range within 10% of the center spans 5-6 grid points, a factor of 4.6-7.9 in LR (0.078-0.36 at L = 1, 0.028-0.22 at L = 40).

Head, norm_bias and whiten_bias never drop by 10% on the smaller side within 20 steps (down to 0.6^20 = 3.7e-5 x the center) at any L. The worst smaller-side loss is 4.10 pp (head, L = 40); most are below 1.5 pp, and at L = 1-2 below 0.5 pp. These three LRs only have an upper limit.

Their larger sides do end, but late and without a pattern across L: head at +7 / +14 / none / none / +3 (L = 1 / 2 / 5 / 10 / 40), norm_bias at none / +17 / +3 / +3 / +6, whiten_bias at +10 / +5 / +6 / +17 / +11. So their Experiment 2 centers mostly mark where "too large" begins, not an optimum: any value below that edge scores about the same.

At L = 10 a few probe points beat the center (whiten_bias +0.57 pp, head +0.15 pp): the joint-search center is not an exact optimum on one seed, which is the center dependence this experiment was meant to remove.

Experiment 5. Which non-conv interval schedules help

Question

The two original 5-interval searches scored 93.77 (conv only) and 93.92 (all four LRs) on held-out seeds, but their conv schedules differ in the last two intervals (0.078, 0.017 vs 0.047, 0.01). With conv fixed, do the other groups' searched intervals help, and which group?

Setup

Item	Value
Schedules	C4 = conv constant 0.13, 0.13, 0.078, 0.047, 0.01 (5 x 40 steps; the four-LR interval search's conv). Searched intervals: head 2100, 2100, 2100, 2100, 1300; norm_bias 770, 770, 770, 460, 280; whiten_bias 460, 460, 280, 280, 280. Baseline: conv 0.22 -> 0, head 1300 -> 0, norm_bias 99 -> 0 (linear over 200 steps), whiten_bias 36 -> 0 over 75 steps, then 0
Runs	Plain training runs (no search) on held-out seeds 100-107, current code; the paired difference is per seed against the first row

Held-out TTA

Conv	head	norm_bias	whiten_bias	Held-out TTA (mean +- SE)	vs first row (paired)
C4	baseline	baseline	baseline	93.95 +- 0.06	-
C4	searched	baseline	baseline	93.90 +- 0.05	-0.05 +- 0.08
C4	baseline	searched	baseline	93.77 +- 0.04	-0.18 +- 0.08
C4	baseline	baseline	searched	93.88 +- 0.05	-0.07 +- 0.07
C4	searched	searched	searched	93.92 +- 0.04	-0.03 +- 0.07
baseline (0.22 -> 0)	searched	searched	searched	93.87 +- 0.04	-0.08 +- 0.07
baseline (0.22 -> 0)	baseline	baseline	baseline	93.92 +- 0.05	-0.03 +- 0.07

For reference, the original conv-only interval search's schedule (conv 0.13, 0.13, 0.078, 0.078, 0.017, the others on the baseline) scored 93.77 on the same seeds (older code).

Observations

The other groups' searched intervals do not help. With conv fixed at C4 and the other three groups on the baseline, the held-out TTA is 93.95 +- 0.06, as high as any row. Switching all three to their searched intervals gives 93.92 (-0.03 +- 0.07 paired); switching one at a time gives head -0.05 +- 0.08, whiten_bias -0.07 +- 0.07 and norm_bias -0.18 +- 0.08 (about 2.3 SE: norm_bias's searched intervals, 770 -> 280 instead of 99 -> 0, probably hurt).

The 93.77 vs 93.92 gap between the original conv-only and four-LR searches came from conv, not from the other groups. Both conv schedules with the other groups on the baseline: 0.13, 0.13, 0.078, 0.078, 0.017 scores 93.77 and C4 (0.13, 0.13, 0.078, 0.047, 0.01) 93.95, +0.18 +- 0.05 paired (7 of 8 seeds). The four-LR rerun reproduces the original run's per-seed scores exactly, so the older 93.77 run is directly comparable.

So 5 constant intervals of conv alone do match the baseline (93.95 vs the baseline's 93.92 on these seeds) when the last two intervals are low enough (0.047, 0.01); Experiment 3's conv-only search stopped at 0.047, 0.028.

The searched non-conv intervals with the baseline's linear conv score 93.87 (-0.08 +- 0.07 vs the C4 control), close to the baseline itself: the other groups' schedules matter little with either conv schedule.

Experiment 6. Seed noise of the baseline

Question

How much does the baseline's score vary across training seeds, and how many seeds does a comparison need?

Setup

Item	Value
Runs	baseline_bs2000, plain training runs on seeds 0-127 (seed 0 is the search seed; 100-107 are the held-out seeds); the seed sets the initialization, data order and augmentation
Score	Test TTA accuracy (%) and test accuracy without TTA

Baseline across seeds

Metric	Seeds	Mean	SD (one run)	SE of the mean	Min	Max
TTA	128	93.942	0.121	0.011	93.61	94.28
No TTA	128	93.164	0.140	0.012	92.69	93.52

Averaging over seeds

Seeds per mean (n)	SD of an n-seed mean: SD/sqrt(n)	Measured (disjoint blocks of n)	Smallest difference two n-seed means resolve (95%)
1	0.121	0.121 (128 blocks)	0.34 pp
2	0.086	0.089 (64 blocks)	0.24 pp
4	0.061	0.069 (32 blocks)	0.17 pp
8	0.043	0.051 (16 blocks)	0.12 pp
16	0.030	0.041 (8 blocks)	0.08 pp
32	0.021	0.028 (4 blocks)	0.06 pp

Seed 0 (the search seed): 93.92. Held-out seeds 100-107: mean 93.92. All seeds: 93.94.

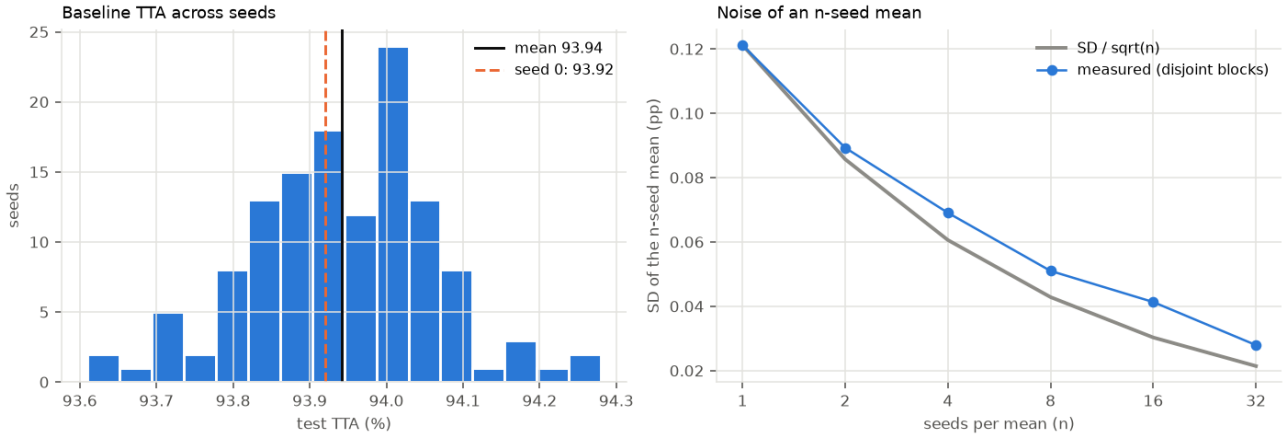


Figure 5. Left: the baseline's TTA over seeds. Right: the spread of an n-seed mean, predicted from the single-run SD and measured on disjoint blocks of seeds.

Observations

One run of the baseline has a seed SD of 0.121 pp in TTA (128 seeds: mean 93.94, range 93.61-94.28, 0.67 pp from worst to best seed); without TTA the SD is 0.140 pp (mean 93.16). The mean over all 128 seeds is known to +-0.011 pp (SE).

Averaging over seeds shrinks the noise about as SD / sqrt(n): measured on disjoint blocks of seeds, an 8-seed mean has SD 0.051 (predicted 0.043; 16 blocks, so the measured value itself is uncertain by about +-20%). Telling two schedules apart at 95% takes a difference of about 0.34 pp with 1 seed each, 0.12 pp with 8 seeds and 0.06 pp with 32 seeds (unpaired; pairing by seed helps only if the seeds' effects are shared).

The seeds used so far are typical: seed 0 (the search seed) scores 93.92 and the held-out seeds 100-107 average 93.92, both within 0.02 pp of the 128-seed mean.

Experiment 7. Hparam range around the baseline

Question

Starting from the baseline instead of a from-scratch center, how far can each LR's initial value move, on each side, before the full run's score drops by 10%?

Setup

Item	Value
Runs	baseline_bs2000, full 200-step runs, seed 0. One LR at a time (the others on the baseline)
Probe	The LR's whole baseline schedule is scaled so its initial value steps along the 0.6^k grid (+1 = x1/0.6, larger), on each side until TTA <= 0.9 x the baseline's TTA on seed 0, or 20 steps. A diverged run counts as a drop

Stop offsets (baseline TTA on seed 0: 93.92; stop at 84.53)

LR	Baseline initial LR	Stop (smaller / larger)	Initial LR range within 10%
conv	0.22	-6 / +6	0.017 - 2.8 (x165)
head	1300	<-20 / +8	0.047 - 46000 (x9.79e+05)
norm_bias	99	<-20 / +9	0.0036 - 5900 (x1.64e+06)
whiten_bias	36	<-20 / >20	0.0013 - 980000 (x7.54e+08)

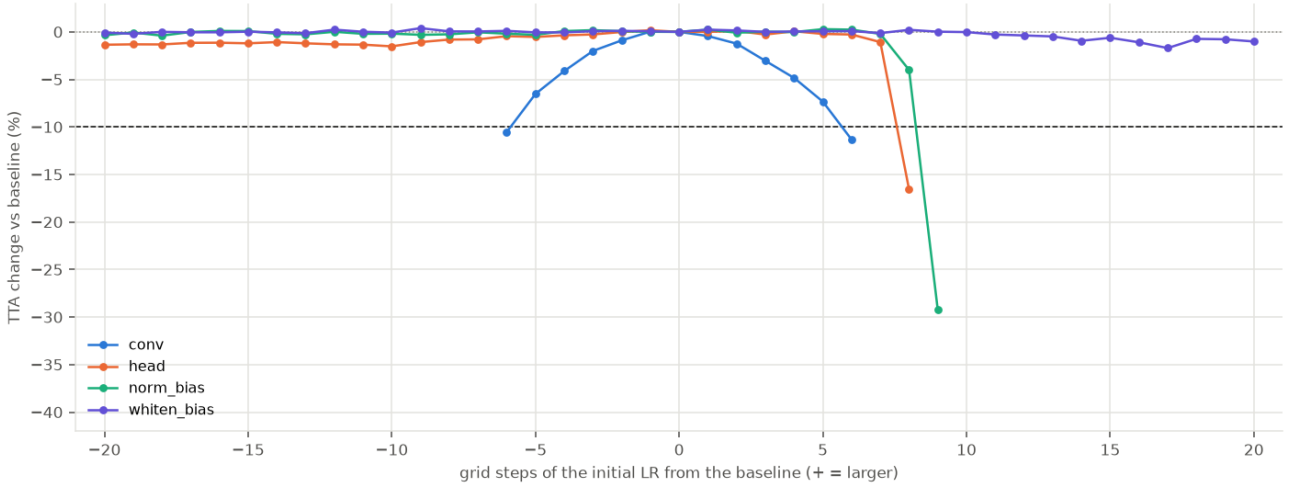


Figure 6. Relative TTA change of full runs as one LR's initial value moves from the baseline; dashed: the 10% stop. Points below -40% are drawn at -40%.

95% confidence set for the best initial LR

Each probed value is one run on seed 0, so its score carries the single-run noise of Experiment 6 (SD 0.121 pp, assumed the same at every value and independent between values). A value is kept when its score is within c of the best observed score, with c = q x SD and q the 95% quantile of max_j (Z_j - Z_1) over the k probed values (Z independent standard normals; Monte Carlo). This is Gupta's subset selection: the set contains the best probed value with probability at least 95%, even in the hardest case where all values are equally good. (open) marks an end that reaches the probe limit on a side without a 10% drop, so the set may continue beyond it.

LR	Probed points k	Cutoff c (pp)	Best observed	95% set (initial LR)	Points in set
conv	13	0.43	0.13 (93.95)	0.13 - 0.22	2 of 13
head	29	0.47	770 (94.08)	280 - 27000	10 of 29
norm_bias	30	0.47	1300 (94.19)	(open) 0.006 - 3500	22 of 30, with gaps
whiten_bias	41	0.48	0.36 (94.29)	(open) 0.0013 - 5900	28 of 41, with gaps

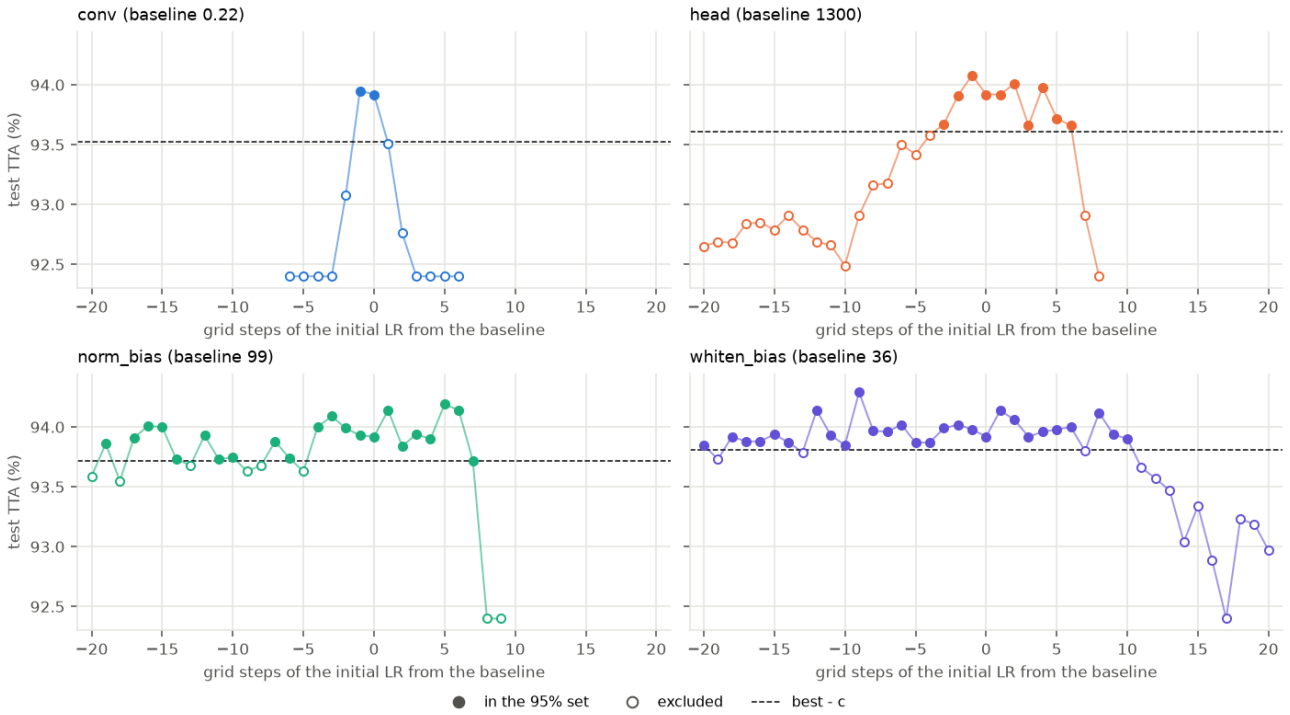


Figure 7. The probes near the top: filled points are in the 95% set (score at least best - c, dashed); points below 92.4 are drawn at 92.4.

Observations

One seed (0); the 10% stop (84.53) is about 80 single-run SDs (Experiment 6) below the baseline, so the stop offsets are not noise; differences under about 0.34 pp between single points are.

Conv has a clear optimum, and the baseline sits on it: 0.13 scores 93.95 and 0.22 (baseline) 93.92; one step further costs 0.4-0.8 pp (0.36: 93.51; 0.078: 93.08). Its 10% range ends at -6 (0.01: 84.03) and +6 (4.6: 83.28), so in full runs conv tolerates a factor of about 165 (0.017 to 2.8) before losing 10%, much wider than in the from-scratch L-step setting of Experiment 4 (5-6 grid points).

Head, norm_bias and whiten_bias never drop by 10% on the smaller side within 20 steps (down to 3.7e-5 x the baseline). Norm_bias and whiten_bias are flat there (all within 93.55-94.29, within about 0.4 pp of the baseline, roughly the single-run noise range); head loses about 1.2 pp once it is below about 8 (92.5-92.9 from 7.7 down to 0.047), still far from 10%.

On the larger side, head stops at +8 (76000: 78.35; 46000 still 92.91) and norm_bias at +9 (9800: 66.44; 5900: 90.16), both after 6-7 steps of no loss: they fail by blowing up, not gradually. whiten_bias never stops: up to 980000 (x27000) it stays above 92.3.

So around the baseline only conv needs tuning; the three other LRs only need to stay below a blow-up edge about 35x above their baseline values (head, norm_bias; the last points without loss are 46000 and 3500), and whiten_bias barely matters at all.

95% confidence sets for the best initial LR (single runs with SD 0.121 pp; cutoff c = 0.43-0.48 pp below the best observed score, wider for more probed points). Conv: 0.13-0.22, the two values at the top (93.95, 93.92); the next ones are excluded (0.36: 93.51, only 0.01 pp under the cutoff 93.52; 0.078: 93.08), so the best conv initial LR is pinned to two grid points, and the baseline's 0.22 is one of them.

Head: 280-27000 (10 of 29 points, contiguous), a factor of about 100 around the best observed 770; below 280 the scores fall off (170: 93.58), and 46000 is 1.0 pp below the baseline (92.91). The baseline's 1300 is inside.

Norm_bias: anything from the smallest probed values (0.006; open, as the probe never dropped 10% on that side) up to 3500, with gaps: the scores below 3500 are within noise of each other, so the data do not locate a best value, only an upper edge. Whiten_bias: likewise from 0.0013 (open) up to 5900; from 9800 upward the scores fall (93.66 and lower) and are excluded.

Experiment 8. Largest LR within the margin, per L (from Experiment 4)

Question

The interval search picks each LR in its first interval as the largest LR scoring within a margin of the best (directional_search with initial_lr_search, initial_dropoff_margin = 0.02). Applied at main interval lengths L = 1, 2, 5, 10, 40 from scratch, does the pick at L = 1 predict the pick at L = 40?

Setup

Item	Value
Data	Experiment 4's probes (no new runs): for each L, one LR at a time on the 0.6^k grid around Experiment 2's center, the other three LRs at the center, TTA right after L steps, seed 0. The probes extend past every point where the rule stops
Rule	As in train_cifar.py: from the center, step up while the score stays within the margin of the best so far, then down the same way (at most 20 steps per side); pick the largest LR within the margin of the best. Margin 0.02 TTA = 2 pp (the code's value); 1 and 0.5 pp shown too

Picks with the code's margin (2 pp): LR (grid steps from the center), and where it lies relative to Experiment 7's full-run 95% range for the best initial LR (last row)

L (steps)	Center TTA	conv	head	norm_bias	whiten_bias
1	25.43	0.36 (+2), above	770 (+5), in	21 (+20)*, in	5900 (+9), in
2	36.89	0.6 (+2), above	460 (+12), in	770 (+12), in	5900 (+3), in
5	52.00	0.13 (+1), in	170 (+20)*, below	5900 (+2), above	9800 (+4), above
10	59.47	0.22 (+2), in	36 (+15), below	2100 (+0), in	2100 (+10), in
40	84.60	0.13 (+1), in	3500 (+1), in	1300 (+3), in	9800 (+3), above
Exp. 7	full run	0.13 - 0.22	280 - 27000	(open) 0.006 - 3500	(open) 0.0013 - 5900

* the upward probe reached the 20-step limit without falling out of the margin, so the pick is the probe limit. in / above / below: inside, above or below the Experiment 7 range ((open): that end of the range is not bounded by the data). Experiment 7 is a different context (full 200-step runs, the other LRs on the baseline schedules), the target a short-horizon pick should predict.

Conv pick by margin

L (steps)	conv, margin 2 pp	conv, margin 1 pp	conv, margin 0.5 pp
1	0.36, above	0.22, in	0.22, in
2	0.6, above	0.36, above	0.22, in
5	0.13, in	0.078, below	0.078, below
10	0.22, in	0.078, below	0.078, below
40	0.13, in	0.078, below	0.078, below
Exp. 7	0.13 - 0.22	0.13 - 0.22	0.13 - 0.22

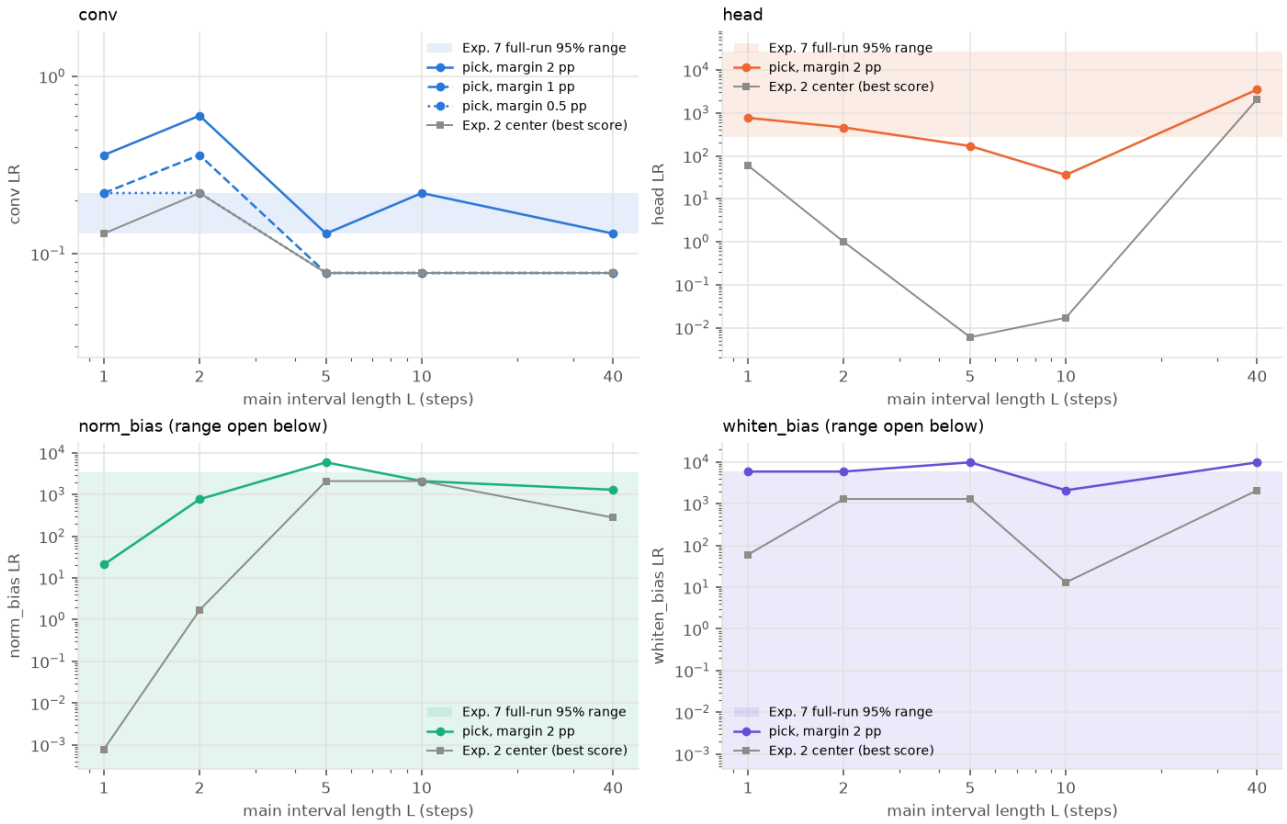


Figure 8. The rule's pick per L (conv at three margins, the others at the code's 2 pp), Experiment 2's best-score center, and the shaded reference: Experiment 7's full-run 95% range for the best initial LR.

Observations

Reference: Experiment 7's full-run 95% ranges for the best initial LR (conv 0.13-0.22, head 280-27000, norm_bias up to 3500, whiten_bias up to 5900; the last two open below, so for them only "above" is a real miss). A good short-horizon pick lands inside.

Conv: with the code's 2 pp margin, L = 1 and L = 2 pick above the range (0.36, 0.6) and L = 5, 10 and 40 inside it (0.13, 0.22, 0.13). So L = 1 does not predict L = 40's pick (2 grid steps apart, 0.36 vs 0.13), and the long-horizon picks are the ones that match the full run. With tighter margins it reverses: at 1 pp and 0.5 pp, L = 1 picks 0.22 (inside) and L = 5-40 pick 0.078 (one step below the range). No single margin puts every L inside.

Why the margin interacts with L: it is absolute, but the score scale is not (the centers score 25.43 at L = 1 and 84.60 at L = 40), and at every L the 2 pp pick sits 1-2 grid steps above that L's best-score center (L = 1: 0.13 -> 0.36; L = 40: 0.078 -> 0.13). The best-score centers alone are inside the range at L = 1 and 2 (0.13, 0.22) and one step below it at L >= 5 (0.078): short horizons favour slightly larger conv LRs than long ones, and the 2 pp rule's upward step compensates at L >= 5 but overshoots at L <= 2.

Head: L = 1 (770), 2 (460) and 40 (3500) pick inside the range; L = 5 (170) and 10 (36) below it, by 1-4 grid steps although the upward probe at L = 5 hit the 20-step limit. norm_bias and whiten_bias: their ranges are open below, so the only misses are picks above the upper edge: norm_bias at L = 5 (5900 > 3500), whiten_bias at L = 5 and 40 (9800 > 5900), each one grid step over. L = 1's picks (21 and 5900) are inside, but their ranges are wide enough that this says little.

So in this data L = 1 lands inside the full-run range for all three of head, norm_bias and whiten_bias (L = 40: two of three), but for conv, the one LR with a narrow range, L = 1 overshoots with the code's margin, while L >= 5 lands inside.

Caveat: in each L the other three LRs sit at that L's own Experiment 2 center, which differs between L (e.g. head 60 at L = 1, 0.006 at L = 5, 2100 at L = 40), so the comparison across L is not in a fixed context, and Experiment 7's context (full runs, baseline schedules) differs from all of them. Running the rule from a common start at every L would remove the first confound; it needs new runs.

Experiment 9. Representation quality of schedule families, noise-aware search

Question

Experiment 3's search scored every candidate on one run (seed 0), so it climbed on seed noise and stopped at schedules that were lucky on seed 0 (e.g. 5 constant intervals: 93.99 on seed 0, 93.62 held-out). With a search that only moves to a neighbour that is statistically better given the seed noise, and adds seeds until it can tell, how close does each family get?

Setup

Item	Value
As Experiment 3	baseline_bs2000, conv only (the others on the baseline); the same families, conditions (no transform, ramp, reverse ramp, wave), grids (LR x0.6, p x0.8) and starts (the oracle projected onto the family); start and final retrained on held-out seeds 100-107
Search	Global neighbour (noisy_family): each round compares the current schedule with all its neighbours (each parameter one grid step up or down; k = 2 for linear, 4 for power and 2 knots, 6 for 3 knots, 10 for 5 intervals and 5 knots). Every point's score is its mean TTA over full runs on search seeds 0, 1, 2, ...
Decision	d = neighbour mean - current mean, SE = 0.12 pp x sqrt(1/n_neighbour + 1/n_current) (the single-run SD of Experiment 6, runs treated as independent). One-sided 80% bounds, Bonferroni over the k neighbours: z = 1.28 / 1.64 / 1.83 / 2.05 for k = 2 / 4 / 6 / 10. Better: d - z SE > 0. Ruled out: d + z SE < 0. Otherwise undecided
Moves and seeds	If any neighbour is better, move to the better one with the highest mean. Else run one more seed on the current schedule and on every undecided neighbour, and repeat. Stop when every neighbour is ruled out, or at 8 seeds per point (undecided neighbours then count as not better)

Held-out TTA of the final schedule (Experiment 3's for comparison; the start; search cost)

Family	No transform	Ramp	Reverse ramp	Wave
5 constant intervals	**93.80** (Exp. 3: 93.62; start 93.52; 2 moves, 119 runs)	**93.68** (Exp. 3: 93.65; start 93.32; 2 moves, 110 runs)	**93.81** (Exp. 3: 93.86; start 93.81; 0 moves, 76 runs)	**93.60** (Exp. 3: 93.82; start 93.11; 1 moves, 77 runs)
Linear decay (1 param)	**93.92** (Exp. 3: 93.86; start 93.92; 0 moves, 15 runs)	**93.42** (Exp. 3: 93.42; start 93.42; 0 moves, 5 runs)	**93.86** (Exp. 3: 93.59; start 93.57; 1 moves, 18 runs)	**93.91** (Exp. 3: 93.91; start 93.72; 1 moves, 8 runs)
Power decay (lr, p)	**93.97** (Exp. 3: 93.86; start 93.92; 1 moves, 36 runs)	**93.70** (Exp. 3: 93.95; start 93.42; 1 moves, 21 runs)	**93.86** (Exp. 3: 93.59; start 93.57; 1 moves, 36 runs)	**93.91** (Exp. 3: 93.91; start 93.72; 1 moves, 22 runs)
2 knots (t = 0, 0.5)	**94.03** (Exp. 3: 93.86; start 93.86; 1 moves, 38 runs)	**93.64** (Exp. 3: 93.77; start 93.14; 1 moves, 29 runs)	**93.85** (Exp. 3: 93.88; start 93.85; 0 moves, 34 runs)	**93.79** (Exp. 3: 93.83; start 93.72; 1 moves, 37 runs)
3 knots (t = 0, 1/3, 2/3)	**93.94** (Exp. 3: 93.86; start 93.94; 0 moves, 48 runs)	**93.88** (Exp. 3: 93.68; start 93.74; 1 moves, 55 runs)	**93.90** (Exp. 3: 93.78; start 93.90; 0 moves, 47 runs)	**94.06** (Exp. 3: 94.06; start 93.90; 1 moves, 60 runs)
5 knots (t = 0, 0.2, ..., 0.8)	**93.97** (Exp. 3: 93.94; start 93.97; 0 moves, 79 runs)	**93.91** (Exp. 3: 93.96; start 93.91; 0 moves, 66 runs)	**93.91** (Exp. 3: 93.91; start 93.91; 0 moves, 78 runs)	**93.97** (Exp. 3: 93.94; start 93.97; 0 moves, 84 runs)

Target: the baseline's 93.92 on the held-out seeds; do-nothing references as in Experiment 3 (ramp 93.42, reverse ramp 93.57, wave 93.72).

Final parameter values (seeds behind the final mean; why the search stopped)

Family	No transform	Ramp	Reverse ramp	Wave
5 constant intervals	0.13, 0.13, 0.13, 0.078, 0.017 (8 seeds; max seeds)	0.6, 0.22, 0.13, 0.028, 0.006 (8 seeds; max seeds)	0.078, 0.078, 0.13, 0.13, 0.078 (8 seeds; max seeds)	0.13, 0.078, 0.13, 0.13, 0.028 (8 seeds; max seeds)
Linear decay (1 param)	0.22 (7 seeds; none can be better)	0.22 (2 seeds; none can be better)	0.13 (8 seeds; none can be better)	0.13 (3 seeds; none can be better)
Power decay (lr, p)	0.22, 1.2 (8 seeds; max seeds)	0.22, 1.2 (7 seeds; none can be better)	0.13, 1 (8 seeds; max seeds)	0.13, 1 (5 seeds; none can be better)
2 knots (t = 0, 0.5)	0.22, 0.078 (8 seeds; none can be better)	1, 0.078 (8 seeds; max seeds)	0.047, 0.13 (8 seeds; max seeds)	0.13, 0.13 (8 seeds; max seeds)
3 knots (t = 0, 1/3, 2/3)	0.22, 0.13, 0.078 (8 seeds; max seeds)	1, 0.22, 0.028 (8 seeds; max seeds)	0.047, 0.078, 0.13 (8 seeds; max seeds)	0.13, 0.078, 0.13 (8 seeds; max seeds)
5 knots (t = 0, 0.2, ..., 0.8)	0.22, 0.22, 0.13, 0.078, 0.047 (8 seeds; max seeds)	1, 0.36, 0.22, 0.078, 0.017 (8 seeds; max seeds)	0.047, 0.078, 0.078, 0.13, 0.13 (8 seeds; max seeds)	0.22, 0.078, 0.078, 0.13, 0.078 (8 seeds; max seeds)

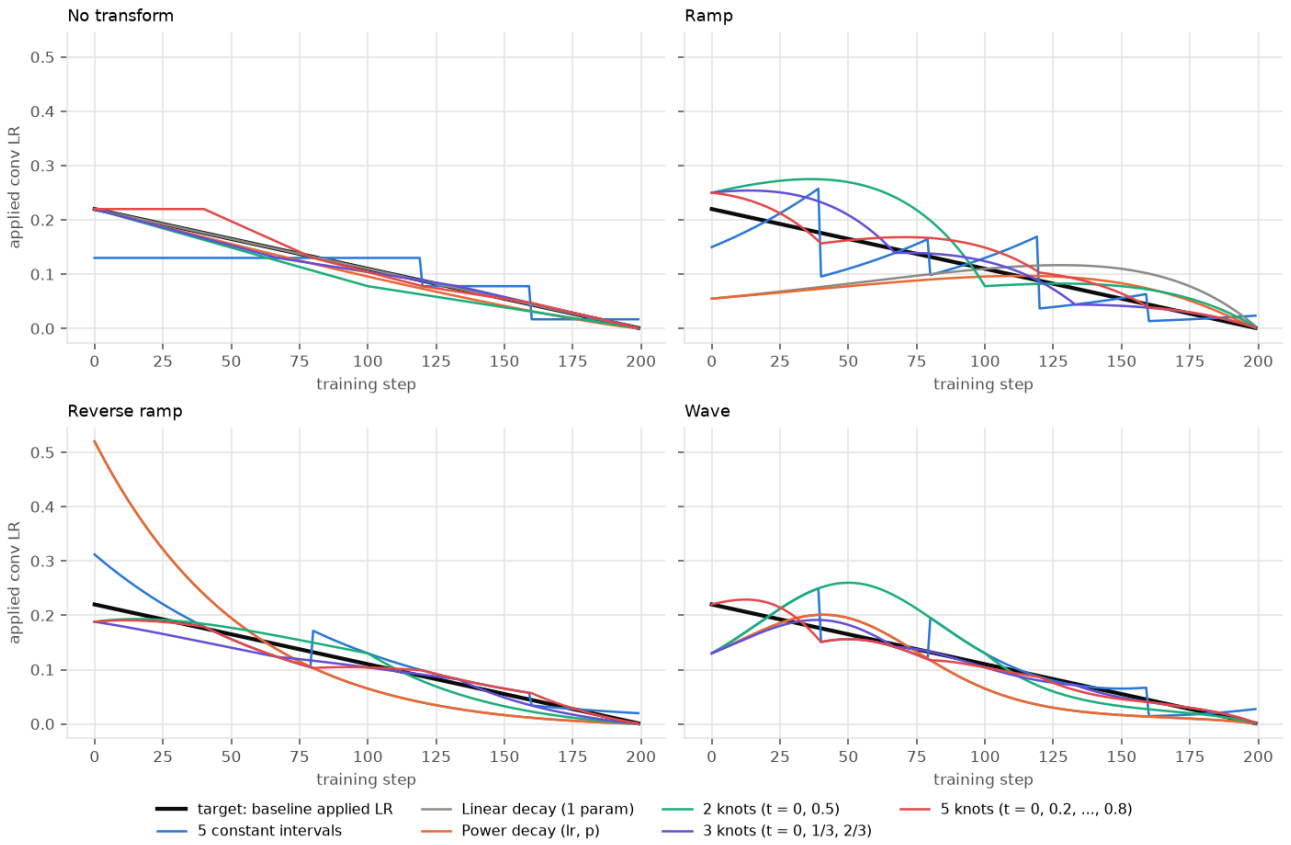


Figure 9. Applied conv LR of each family's final schedule (noise-aware search); the black line is the baseline's applied LR.

Observations

Held-out means over 8 seeds have a standard error of about 0.045 pp, so differences below about 0.1 pp between cells are not resolved. Target: the baseline's 93.92.

The search's own score is now honest. Its final score (the mean over the seeds the search ran it on) is 0.03 pp below the held-out score on average (SD 0.06 over the 24 searches); in Experiment 3 the seed-0 score was 0.15 pp above it (SD 0.15), the winner's curse. Example, 5 constant intervals without a transform: search 93.77, held-out 93.80 (Experiment 3: 93.99 on seed 0, 93.62 held-out).

It does not make the start worse: no final schedule is worse than its start by more than 0.05 pp held-out (mean gain +0.14 pp over the 24 searches); in Experiment 3, 5 of 24 finals were worse than their start by more than 0.05 pp.

It is conservative. It made 0-2 moves per search (Experiment 3: 0-6), and 17 of 24 searches stopped at the 8-seed cap with undecided neighbours: with 8 seeds, a move needs a measured gain of about 0.08-0.12 pp (z = 1.28-2.05 times 0.06 pp), so smaller true gains are left on the table. Against Experiment 3 the held-out results are +0.03 pp

on average: better by more than 0.1 pp in 7 cells (e.g. 5 intervals 93.80 vs 93.62, linear and power under the reverse ramp 93.86 vs 93.59, 2 knots 94.03 vs 93.86), worse in 3 (power under the ramp 93.70 vs 93.95, 5 intervals under the wave 93.60 vs 93.82, 2 knots under the ramp 93.64 vs 93.77), where Experiment 3's longer single-seed walk happened to reach a better region (power, ramp: lr 0.36, p = 2, one and three grid steps from the start).

5 constant intervals without a transform reach 93.80, still below C4's 93.95 (Experiment 5): the search stopped at 0.13, 0.13, 0.13, 0.078, 0.017, three grid steps from C4 (0.13, 0.13, 0.078, 0.047, 0.01), with neighbours still undecided at 8 seeds.

Families: 5 knots is at the target in every condition (93.91-93.97) without moving, as its projected start already was. Linear decay cannot compensate the ramp (93.42, the do-nothing score), as in Experiment 3. The best family per condition: 2 knots 94.03 (no transform), 5 knots 93.91 (ramp, reverse ramp), 3 knots 94.06 (wave).

Cost: 1198 training runs for the 24 searches (Experiment 3: 270), 524 s per search on average (Experiment 3: 114 s), plus 16 held-out runs each.

Experiment 10. Representation quality of schedule families, guided search

Question

Experiment 3's single-seed search walked far but climbed on seed noise; Experiment 9's significance test removed the noise but rarely moved more than one grid step. Can a search that explores deeper on cheap single runs, pools them with a smooth fit, jumps several steps (and several parameters) at once, and accepts a move only on fresh seeds, get both?

Setup

Item	Value
As Experiment 3	baseline_bs2000, conv only (the others on the baseline); the same families, conditions, grids (LR x0.6, p x0.8) and starts; start and final retrained on held-out seeds 100-107
1. Explore	From the current schedule, along each parameter and direction, one grid step at a time, one run each on seed 0, while the run is not clearly worse than the best run on that line (score >= best - 0.84 x 0.12 pp x sqrt(2), one-sided 80%), up to 3 steps
2. Fit	Per parameter, a least-squares quadratic in the grid offset through that line's seed-0 runs (the current schedule's included): the smoothness prior pools 4-7 runs. Predicted gain of an offset = fit there - fit at the current schedule
3. Propose	Each parameter moved alone to its best fitted offset, and the joint move of every parameter with a positive predicted gain (gains summed, a separability prior); the 3 largest predicted gains
4. Verify	On fresh seeds 1-8 (never seed 0, which chose the candidates): better if the lower one-sided 80% bound of (candidate - current) > 0, ruled out if the upper bound < 0, Bonferroni over the candidates (z = 1.50 for 3), SE = 0.12 pp x sqrt(1/n_c + 1/n_current); otherwise one more seed for the current schedule and the undecided candidates, up to 8
Moves	Move to the better candidate with the highest mean and repeat; stop when no parameter is predicted to gain or no candidate proves better

Held-out TTA of the final schedule (Experiments 3 and 9 for comparison; the start; search cost)

Family	No transform	Ramp	Reverse ramp	Wave
5 constant intervals	**93.52** (Exp. 3: 93.62; Exp. 9: 93.80; start 93.52; 0 moves, 53 runs)	**93.78** (Exp. 3: 93.65; Exp. 9: 93.68; start 93.32; 2 moves, 91 runs)	**93.73** (Exp. 3: 93.86; Exp. 9: 93.81; start 93.81; 2 moves, 97 runs)	**93.84** (Exp. 3: 93.82; Exp. 9: 93.60; start 93.11; 2 moves, 100 runs)
Linear decay (1 param)	**93.92** (Exp. 3: 93.86; Exp. 9: 93.92; start 93.92; 0 moves, 4 runs)	**93.42** (Exp. 3: 93.42; Exp. 9: 93.42; start 93.42; 0 moves, 4 runs)	**93.86** (Exp. 3: 93.59; Exp. 9: 93.86; start 93.57; 1 moves, 7 runs)	**93.91** (Exp. 3: 93.91; Exp. 9: 93.91; start 93.72; 1 moves, 8 runs)
Power decay (lr, p)	**93.92** (Exp. 3: 93.86; Exp. 9: 93.97; start 93.92; 0 moves, 9 runs)	**93.88** (Exp. 3: 93.95; Exp. 9: 93.70; start 93.42; 2 moves, 29 runs)	**93.85** (Exp. 3: 93.59; Exp. 9: 93.86; start 93.57; 2 moves, 30 runs)	**93.91** (Exp. 3: 93.91; Exp. 9: 93.91; start 93.72; 1 moves, 15 runs)
2 knots (t = 0, 0.5)	**93.86** (Exp. 3: 93.86; Exp. 9: 94.03; start 93.86; 0 moves, 5 runs)	**93.86** (Exp. 3: 93.77; Exp. 9: 93.64; start 93.14; 2 moves, 20 runs)	**93.85** (Exp. 3: 93.88; Exp. 9: 93.85; start 93.85; 0 moves, 16 runs)	**93.95** (Exp. 3: 93.83; Exp. 9: 93.79; start 93.72; 2 moves, 32 runs)
3 knots (t = 0, 1/3, 2/3)	**93.94** (Exp. 3: 93.86; Exp. 9: 93.94; start 93.94; 0 moves, 22 runs)	**93.88** (Exp. 3: 93.68; Exp. 9: 93.88; start 93.74; 1 moves, 25 runs)	**93.94** (Exp. 3: 93.78; Exp. 9: 93.90; start 93.90; 1 moves, 55 runs)	**94.06** (Exp. 3: 94.06; Exp. 9: 94.06; start 93.90; 1 moves, 34 runs)
5 knots (t = 0, 0.2, ..., 0.8)	**93.97** (Exp. 3: 93.94; Exp. 9: 93.97; start 93.97; 0 moves, 30 runs)	**93.91** (Exp. 3: 93.96; Exp. 9: 93.91; start 93.91; 0 moves, 50 runs)	**93.91** (Exp. 3: 93.91; Exp. 9: 93.91; start 93.91; 0 moves, 11 runs)	**93.97** (Exp. 3: 93.94; Exp. 9: 93.97; start 93.97; 0 moves, 22 runs)

Target: the baseline's 93.92 on the held-out seeds; do-nothing references as in Experiment 3 (ramp 93.42, reverse ramp 93.57, wave 93.72).

Final parameter values (fresh seeds behind the final mean, 'not verified' when the search stopped before testing any candidate; why it stopped)

Family	No transform	Ramp	Reverse ramp	Wave
5 constant intervals	0.22, 0.13, 0.13, 0.078, 0.028 (8 seeds; none proved better)	0.36, 0.22, 0.078, 0.028, 0.0036 (8 seeds; none proved better)	0.047, 0.078, 0.13, 0.13, 0.047 (8 seeds; none proved better)	0.078, 0.078, 0.13, 0.078, 0.01 (8 seeds; none proved better)
Linear decay (1 param)	0.22 (not verified; no predicted gain)	0.22 (not verified; no predicted gain)	0.13 (1 seeds; no predicted gain)	0.13 (2 seeds; no predicted gain)
Power decay (lr, p)	0.22, 1 (not verified; no predicted gain)	0.36, 1.6 (2 seeds; none proved better)	0.13, 0.64 (3 seeds; none proved better)	0.13, 1 (2 seeds; no predicted gain)
2 knots (t = 0, 0.5)	0.22, 0.13 (not verified; no predicted gain)	0.6, 0.047 (1 seeds; no predicted gain)	0.047, 0.13 (3 seeds; none proved better)	0.017, 0.13 (3 seeds; no predicted gain)
3 knots (t = 0, 1/3, 2/3)	0.22, 0.13, 0.078 (4 seeds; none proved better)	1, 0.22, 0.028 (1 seeds; no predicted gain)	0.047, 0.078, 0.078 (8 seeds; none proved better)	0.13, 0.078, 0.13 (4 seeds; no predicted gain)
5 knots (t = 0, 0.2, ..., 0.8)	0.22, 0.22, 0.13, 0.078, 0.047 (5 seeds; none proved better)	1, 0.36, 0.22, 0.078, 0.017 (8 seeds; none proved better)	0.047, 0.078, 0.078, 0.13, 0.13 (not verified; no predicted gain)	0.22, 0.078, 0.078, 0.13, 0.078 (3 seeds; none proved better)

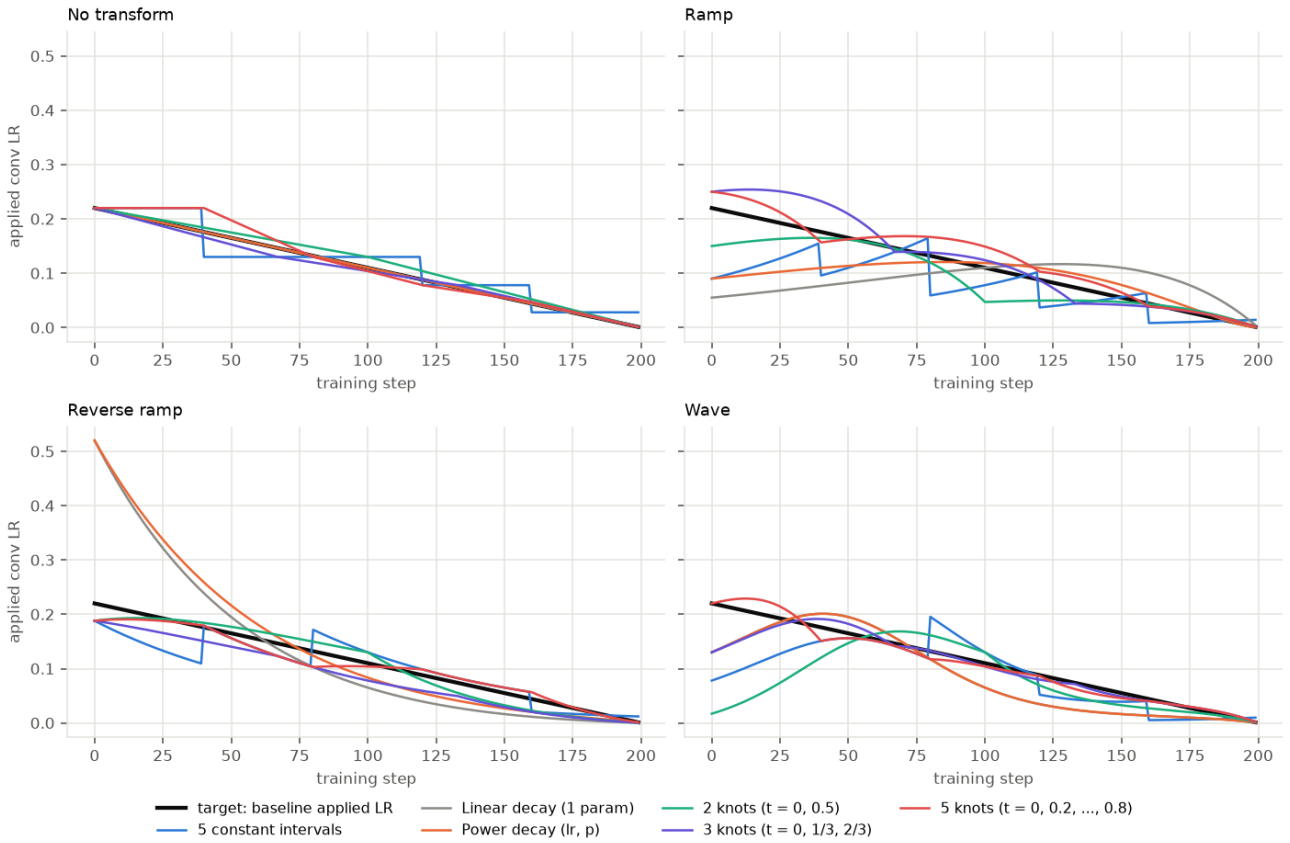


Figure 10. Applied conv LR of each family's final schedule (guided search); the black line is the baseline's applied LR.

Observations

Held-out means over 8 seeds have a standard error of about 0.045 pp, so differences below about 0.1 pp between cells are not resolved. Target: the baseline's 93.92.

Overall the guided search is slightly ahead at lower cost, but not decisively. Mean held-out TTA over the 24 cells: 93.86 (Experiment 9: 93.85, Experiment 3: 93.81); 14 cells within 0.05 pp of the target or above (Experiment 9: 13, Experiment 3: 9). Against Experiment 9 it is better by more than 0.1 pp in 4 cells (5 intervals under the wave 93.84 vs 93.60, power under the ramp 93.88 vs 93.70, 2 knots under the ramp 93.86 vs 93.64 and under the wave 93.95 vs 93.79) and worse in 2 (5 intervals without a transform 93.52 vs 93.80, 2 knots without a transform 93.86 vs 94.03). Cost: 769 training runs for the 24 searches (Experiment 9: 1198, Experiment 3: 270), 337 s per search on average (524 s, 114 s).

Deeper moves happen. Under the ramp, power reached lr 0.36, p = 1.6 (93.88; Experiment 3's lucky walk found lr 0.36, p = 2 at 93.95, Experiment 9 stopped at p = 1.2 with 93.70), and 2 knots reached 0.6, 0.047 (93.86 vs 93.64); 20 accepted moves in total (Experiment 9: 16), several of them multi-step or joint.

The fresh-seed gate keeps the search honest: the final schedule's search score is 0.02 pp above its held-out score on average (SD 0.075, over the 19 searches that verified their final; Experiment 9: -0.03, Experiment 3: +0.15), and only one final is worse than its start by more than 0.05 pp (5 intervals under the reverse ramp, 93.73 vs 93.81).

The weak point is candidate selection. The seed-0 runs along each line carry the same 0.12 pp noise as the effects being measured, so the quadratic fits often follow noise to the end of the explored range. 5 intervals without a transform shows it: the fit sent v0 three steps down to 0.047 on one lucky run (+0.14 pp predicted), the joint move summed four such noisy gains (+0.48 pp predicted), and with v4 to 0.01 these filled all 3 candidate slots; on fresh seeds they measured -0.16, +0.02 and -0.08 pp, so the search stopped at its start (93.52). The first step toward C4, v2 one step down to 0.078, was predicted only +0.05 pp and never tested (v3's line predicted no gain).

Five searches stopped in their first round without testing any candidate (no parameter predicted to gain; 'not verified' in the table); their final is the start.

Likely fixes, not run: rank candidates by predicted gain relative to the fit's own standard error rather than by the raw gain, build the joint move only from parameters whose fitted gain is clearly positive, and test the single-parameter moves before the joint one.

Experiment 11. Representation quality of schedule families, guided search with uncertainty-aware selection

Question

Experiment 10's weak point was candidate selection: quadratic fits through noisy single runs followed noise to the end of the explored range, and a joint move summing those noisy gains crowded out the real improvements. Does ranking by the fit's lower confidence bound, building the joint move only from clearly positive parameters, and testing single moves first fix it?

Setup

Item	Value
As Experiment 10	Same families, conditions, starts, exploration (one seed-0 run per point, up to 3 steps per direction), quadratic fit per parameter, fresh-seed verification (seeds 1-8, one-sided 80%, Bonferroni over the candidates raced together) and held-out evaluation
Fit uncertainty	The SE of each fitted gain (fit at the offset - fit at the current schedule) from the regression with 0.12 pp noise per run: large where the fit extrapolates or rests on few points
Offset choice	Per parameter, the offset (other than the current one) with the highest lower bound, gain - 0.84 SE
Candidates	Single-parameter moves with a positive fitted gain, ranked by that lower bound; the best 3 are raced first. Only if none proves better: the joint move of the parameters whose lower bound is above 0 (at least 2), raced alone

Held-out TTA of the final schedule (Experiments 9 and 10 for comparison; the start; search cost)

Family	No transform	Ramp	Reverse ramp	Wave
5 constant intervals	**93.52** (Exp. 9: 93.80; Exp. 10: 93.52; start 93.52; 0 moves, 63 runs)	**93.74** (Exp. 9: 93.68; Exp. 10: 93.78; start 93.32; 2 moves, 100 runs)	**93.73** (Exp. 9: 93.81; Exp. 10: 93.73; start 93.81; 2 moves, 94 runs)	**93.66** (Exp. 9: 93.60; Exp. 10: 93.84; start 93.11; 1 moves, 47 runs)
Linear decay (1 param)	**93.92** (Exp. 9: 93.92; Exp. 10: 93.92; start 93.92; 0 moves, 4 runs)	**93.42** (Exp. 9: 93.42; Exp. 10: 93.42; start 93.42; 0 moves, 4 runs)	**93.86** (Exp. 9: 93.86; Exp. 10: 93.86; start 93.57; 1 moves, 7 runs)	**93.91** (Exp. 9: 93.91; Exp. 10: 93.91; start 93.72; 1 moves, 8 runs)
Power decay (lr, p)	**93.92** (Exp. 9: 93.97; Exp. 10: 93.92; start 93.92; 0 moves, 9 runs)	**93.88** (Exp. 9: 93.70; Exp. 10: 93.88; start 93.42; 2 moves, 29 runs)	**93.85** (Exp. 9: 93.86; Exp. 10: 93.85; start 93.57; 3 moves, 44 runs)	**93.91** (Exp. 9: 93.91; Exp. 10: 93.91; start 93.72; 1 moves, 15 runs)
2 knots (t = 0, 0.5)	**93.86** (Exp. 9: 94.03; Exp. 10: 93.86; start 93.86; 0 moves, 5 runs)	**93.86** (Exp. 9: 93.64; Exp. 10: 93.86; start 93.14; 2 moves, 19 runs)	**93.85** (Exp. 9: 93.85; Exp. 10: 93.85; start 93.85; 0 moves, 16 runs)	running
3 knots (t = 0, 1/3, 2/3)	**93.94** (Exp. 9: 93.94; Exp. 10: 93.94; start 93.94; 0 moves, 22 runs)	**93.88** (Exp. 9: 93.88; Exp. 10: 93.88; start 93.74; 1 moves, 25 runs)	**93.94** (Exp. 9: 93.90; Exp. 10: 93.94; start 93.90; 1 moves, 50 runs)	running
5 knots (t = 0, 0.2, ..., 0.8)	**93.97** (Exp. 9: 93.97; Exp. 10: 93.97; start 93.97; 0 moves, 30 runs)	**93.93** (Exp. 9: 93.91; Exp. 10: 93.91; start 93.91; 2 moves, 130 runs)	**93.91** (Exp. 9: 93.91; Exp. 10: 93.91; start 93.91; 0 moves, 11 runs)	running

Target: the baseline's 93.92 on the held-out seeds; do-nothing references as in Experiment 3 (ramp 93.42, reverse ramp 93.57, wave 93.72).

Final parameter values (fresh seeds behind the final mean, 'not verified' when the search stopped before testing any candidate; why it stopped)

Family	No transform	Ramp	Reverse ramp	Wave
5 constant intervals	0.22, 0.13, 0.13, 0.078, 0.028 (8 seeds; none proved better)	0.6, 0.22, 0.13, 0.028, 0.0036 (8 seeds; none proved better)	0.047, 0.078, 0.13, 0.13, 0.047 (8 seeds; none proved better)	0.13, 0.078, 0.13, 0.13, 0.017 (5 seeds; none proved better)
Linear decay (1 param)	0.22 (not verified; no predicted gain)	0.22 (not verified; no predicted gain)	0.13 (1 seeds; no predicted gain)	0.13 (2 seeds; no predicted gain)
Power decay (lr, p)	0.22, 1 (not verified; no predicted gain)	0.36, 1.6 (2 seeds; none proved better)	0.13, 0.64 (3 seeds; none proved better)	0.13, 1 (2 seeds; no predicted gain)
2 knots (t = 0, 0.5)	0.22, 0.13 (not verified; no predicted gain)	0.6, 0.047 (1 seeds; no predicted gain)	0.047, 0.13 (3 seeds; none proved better)	running
3 knots (t = 0, 1/3, 2/3)	0.22, 0.13, 0.078 (4 seeds; none proved better)	1, 0.22, 0.028 (1 seeds; no predicted gain)	0.047, 0.078, 0.078 (8 seeds; none proved better)	running
5 knots (t = 0, 0.2, ..., 0.8)	0.22, 0.22, 0.13, 0.078, 0.047 (5 seeds; none proved better)	0.6, 0.36, 0.22, 0.047, 0.01 (8 seeds; none proved better)	0.047, 0.078, 0.078, 0.13, 0.13 (not verified; no predicted gain)	running

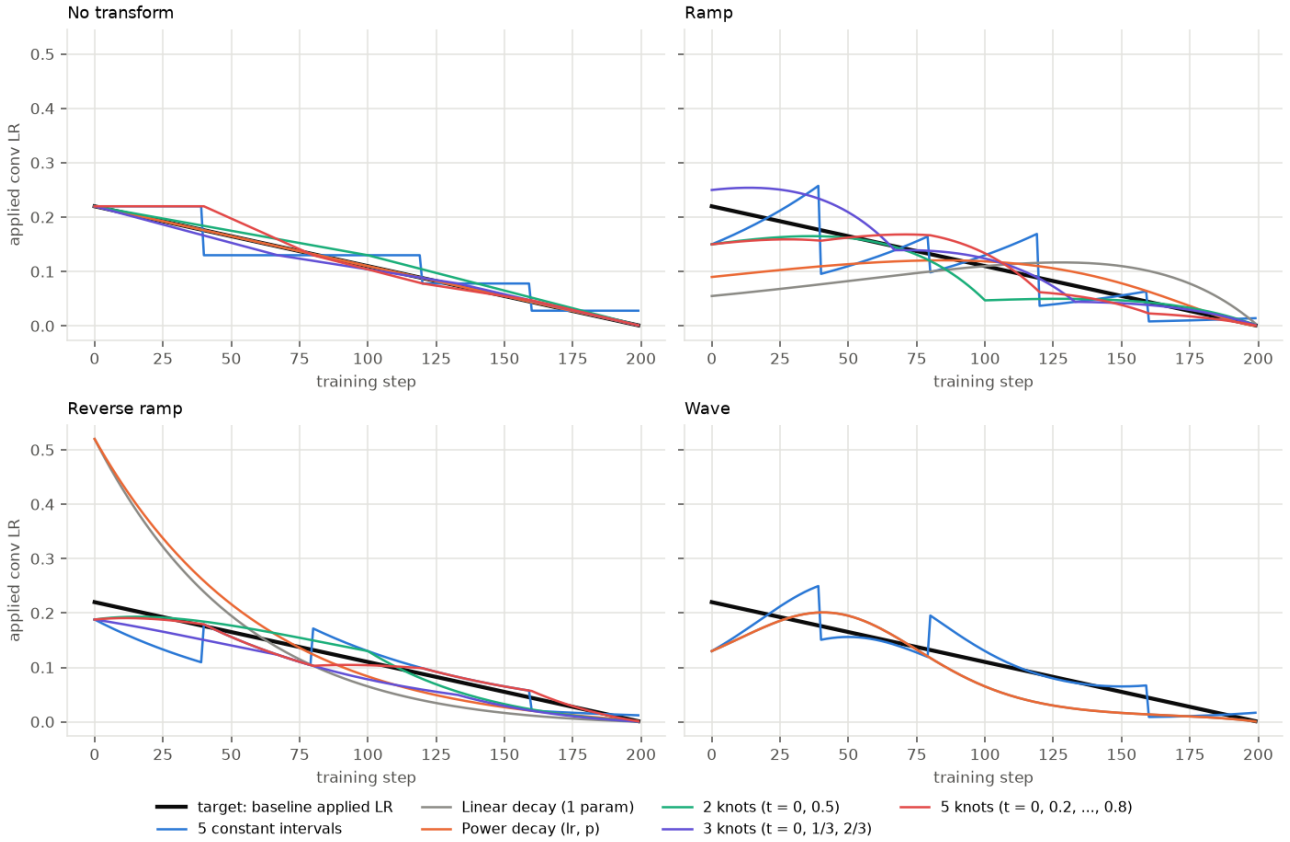


Figure 11. Applied conv LR of each family's final schedule (guided search, uncertainty-aware selection); the black line is the baseline's applied LR.

Experiment 12. Which lookahead decides constant intervals like the final result, and how noisy is it?

Question

An interval search commits each constant interval by scoring candidates with a lookahead from the same trained state. Which lookahead ranks the candidates like their true effect on the final score, and how noisy is one seed's comparison of two candidates scored from the same state?

Setup

Item	Value
Plan	BS 2000, conv as 5 constant intervals of 40 steps following C4 (0.13, 0.13, 0.078, 0.047, 0.01; it matches the baseline, Experiment 5), the other groups on the baseline
Candidates	At interval k, from the plan's trained state: the plan's value moved -2..+2 grid steps (x0.6), trained for the interval
Scores	truth: the rest of the run on the plan (what the decision affects); rest_lin: the rest of the run with conv decaying linearly from the candidate to 0; c40_m: one interval (40 steps) at m x candidate (m = 0.6, 0.36, 0.22); l40: one interval decaying linearly to 0
Seeds	8 (0-7); each seed is its own run from initialization, so the plan's state differs per seed
Decision quality	1-seed hit: how often one seed's best candidate by the score is the candidate with the best 8-seed mean truth; regret: the mean, over seeds, of the mean truth lost by that pick (pp)

8-seed pick; 1-seed hit rate; 1-seed regret (pp of final TTA)

Interval (best by truth)	truth	rest_lin	c40_0.22	c40_0.36	c40_0.6	l40
1 (0.22)	0.22; 0.25; 0.039	0.13; 0.38; 0.001	0.22; 0.38; 0.001	0.13; 0.00; 0.001	0.078; 0.00; 0.032	0.13; 0.00; 0.001
2 (0.13)	0.13; 0.50; 0.028	0.13; 0.88; 0.011	0.078; 0.38; 0.082	0.078; 0.00; 0.115	0.047; 0.00; 0.294	0.078; 0.12; 0.079
3 (0.078)	0.078; 0.62; 0.009	0.078; 0.62; 0.017	0.078; 0.50; 0.046	0.047; 0.12; 0.080	0.028; 0.00; 0.134	0.047; 0.25; 0.068
4 (0.047)	0.047; 0.50; 0.029	0.047; 0.25; 0.050	0.047; 0.50; 0.029	0.028; 0.25; 0.037	0.028; 0.00; 0.083	0.047; 0.25; 0.050
5 (0.01)	0.01; 0.50; 0.027	-	-	-	-	-

Noise of one seed's comparison of two candidates

Interval	Noise of a paired comparison (SD over seeds of candidate - plan value, pp): truth / rest_lin
1	0.047: 0.23 / 0.17; 0.078: 0.27 / 0.22; 0.22: 0.24 / 0.15; 0.36: 0.25 / 0.19
2	0.047: 0.20 / 0.17; 0.078: 0.12 / 0.16; 0.22: 0.21 / 0.14; 0.36: 0.16 / 0.25
3	0.028: 0.16 / 0.19; 0.047: 0.14 / 0.19; 0.13: 0.15 / 0.22; 0.22: 0.20 / 0.16
4	0.017: 0.10 / 0.09; 0.028: 0.10 / 0.07; 0.078: 0.17 / 0.13; 0.13: 0.15 / 0.19
5	0.0036: 0.07 / -; 0.006: 0.06 / -; 0.017: 0.10 / -; 0.028: 0.13 / -

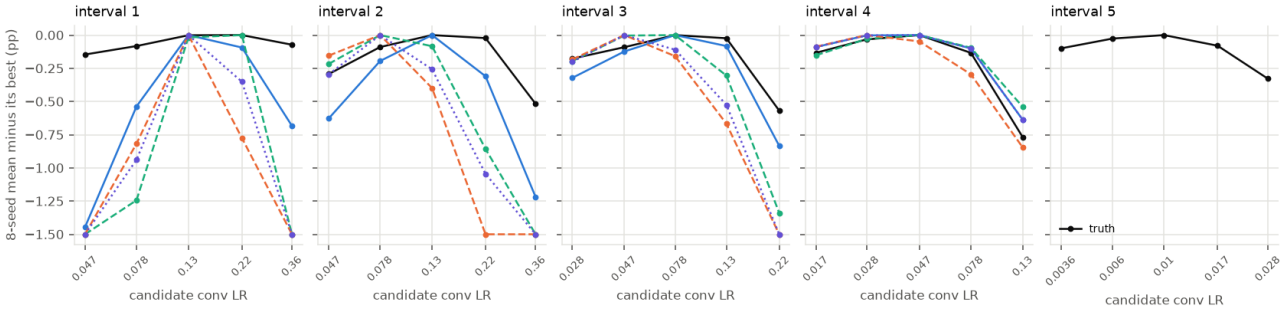


Figure 12. Each score's 8-seed mean per candidate, relative to its own best (points below -1.5 drawn at -1.5). A lookahead decides like the truth when its curve peaks where the truth's does.

Observations

The rest-of-run linear lookahead decides like the truth: its 8-seed pick is the truth's best in intervals 2-4 (interval 5 has no lookahead), and in interval 1 it picks 0.13 where the truth ties 0.13 and 0.22 (93.83 both). One seed's pick by it loses 0.001-0.050 pp of final TTA per interval, about what one seed's pick by the truth itself loses (0.009-0.039): the truth is flat near its optimum, so a neighbouring pick costs little.

One-interval (40-step) lookaheads are biased toward small LRs once training is under way: in interval 2 every one of them picks 0.047-0.078 where the truth peaks at 0.13 (regret 0.08-0.29 pp per seed), and c40_0.6 picks 1-2 grid steps below the truth's best in intervals 2-4. The mildest cooldown, c40_0.22, is the least biased (regret 0.03-0.08 pp from interval 2 on) but still loses to rest_lin in intervals 2 and 3.

Scoring candidates from the same trained state does not cancel the seed noise: one seed's difference between two candidates' truths has an SD of 0.12-0.27 pp in intervals 1-2, as large as two independent full runs (0.17), falling to 0.06-0.17 pp in intervals 4-5 as fewer steps remain. Differences of 0.05-0.1 pp between neighbouring candidates are therefore not resolved by one seed, but they also cost little.

Cost at BS 2000: rest_lin costs the rest of the run per candidate (200, 160, 120, 80, 40 steps for intervals 1-5), so about 3 candidates per interval cost 1800 steps = 9 x a run for one searched LR, within the campaign's 20x budget.

Experiment 13. ICS v1: constant intervals with a rest-of-run lookahead, all batch sizes and transforms

Question

Does the interval-cooldown search, with the lookahead Experiment 12 found faithful, match the baseline at batch sizes 2000, 500 and 125, with and without transforms, searching conv only or all four LRs, within 20 full runs?

Setup

Item	Value
Representation	Each searched LR: 5 constant intervals (spec values; the transform applies on top), on the x0.6 grid
Per interval	Climb the lead LR (conv) from the previous interval's value, smaller values first, moving while the score improves. A candidate's score: train the interval from the committed state, then the rest of the run with every searched LR decaying linearly from its candidate value to 0 (rest_lin); the last interval is scored by the final TTA. Commit the best and continue from its trained state
Conv only (a)	conv from its baseline initial LR; head, norm_bias, whiten_bias on the baseline
4 LRs (b)	All four from 1.0. Interval 1: each LR's level is scanned on a x7.7 grid (both neighbours first, then the better side while improving), scored after the interval plus one interval at 0.36 x the value; then conv is refined on the x0.6 grid with rest_lin. Later intervals: conv is climbed and the other three follow its profile (value_k = value_1 x conv_k / conv_1)
Budget	At most 20 full runs of train steps (every trained step counts); the transform applies to the searched groups
Evaluation	2 search seeds (0, 1) per cell; each committed schedule retrained on held-out seeds 100-107

Held-out TTA (mean over search seeds; vs the baseline; per search seed); cost in full runs

Batch size, scenario (baseline)	none	ramp	reverse ramp	wave
BS 2000, conv only (93.92)	**93.82** (-0.10; 93.82/93.82); 10.0x	**93.15** (-0.77; 93.19/93.11); 10.6x	**93.80** (-0.12; 93.81/93.79); 11.0x	**93.54** (-0.38; 93.61/93.46); 9.6x
BS 2000, 4 LRs (93.92)	**92.85** (-1.07; 93.68/92.01); 17.2x	**91.65** (-2.27; 90.82/92.48); 15.6x	**93.70** (-0.22; 93.89/93.51); 15.2x	**92.17** (-1.75; 91.94/92.40); 15.6x
BS 500, conv only (94.00)	**93.84** (-0.16; 93.75/93.92); 9.4x	**93.40** (-0.60; 93.27/93.53); 10.4x	**93.88** (-0.11; 93.95/93.82); 11.8x	**93.50** (-0.50; 93.25/93.75); 10.4x
BS 500, 4 LRs (94.00)	**93.58** (-0.42; 93.53/93.62); 15.8x	**92.93** (-1.07; 92.79/93.07); 16.6x	running	running
BS 125, conv only (93.86)	**93.73** (-0.13; 93.70/93.76); 9.6x	running	**93.66** (-0.20; 93.72/93.60); 10.2x	running
BS 125, 4 LRs (93.86)	running	running	running	running

Cells within 0.1 pp of the baseline (about 2 SE of an 8-seed mean) or above: 1 of 16 finished. * = the search hit its step budget. Do-nothing: the baseline schedule kept under the transform.

Committed schedules (search seed 0)

Cell	Committed schedules (search seed 0)
BS 2000, conv only, none	conv: 0.13, 0.13, 0.13, 0.047, 0.006
BS 2000, conv only, ramp	conv: 0.13, 0.13, 0.047, 0.017, 0.0036
BS 2000, conv only, reverse ramp	conv: 0.078, 0.13, 0.078, 0.078, 0.028
BS 2000, conv only, wave	conv: 0.22, 0.13, 0.13, 0.078, 0.01
BS 2000, 4 LRs, none	conv: 0.13, 0.13, 0.047, 0.028, 0.006; head: 3500, 3500, 1300, 750, 160; norm_bias: 0.13, 0.13, 0.047, 0.028, 0.006; whiten_bias: 460, 460, 170, 99, 21
BS 2000, 4 LRs, ramp	conv: 0.13, 0.13, 0.047, 0.017, 0.0022; head: 1, 1, 0.36, 0.13, 0.017; norm_bias: 1, 1, 0.36, 0.13, 0.017; whiten_bias: 0.13, 0.13, 0.047, 0.017, 0.0022
BS 2000, 4 LRs, reverse ramp	conv: 0.078, 0.078, 0.078, 0.078, 0.028; head: 460, 460, 460, 460, 170; norm_bias: 60, 60, 60, 60, 22; whiten_bias: 1, 1, 1, 1, 0.36
BS 2000, 4 LRs, wave	conv: 0.13, 0.13, 0.078, 0.078, 0.01; head: 1, 1, 0.6, 0.6, 0.077; norm_bias: 1, 1, 0.6, 0.6, 0.077; whiten_bias: 1, 1, 0.6, 0.6, 0.077
BS 500, conv only, none	conv: 0.078, 0.078, 0.047, 0.047, 0.01
BS 500, conv only, ramp	conv: 0.078, 0.078, 0.028, 0.01, 0.0022
BS 500, conv only, reverse ramp	conv: 0.028, 0.047, 0.047, 0.028, 0.017
BS 500, conv only, wave	conv: 0.13, 0.13, 0.078, 0.047, 0.01
BS 500, 4 LRs, none	conv: 0.078, 0.047, 0.047, 0.028, 0.006; head: 460, 280, 280, 170, 35; norm_bias: 0.017, 0.01, 0.01, 0.0061, 0.0013; whiten_bias: 7.7, 4.6, 4.6, 2.8, 0.59
BS 500, 4 LRs, ramp	conv: 0.078, 0.047, 0.028, 0.01, 0.0013; head: 3500, 2100, 1300, 450, 58; norm_bias: 7.7, 4.6, 2.8, 0.99, 0.13; whiten_bias: 1, 0.6, 0.36, 0.13, 0.017
BS 125, conv only, none	conv: 0.028, 0.028, 0.017, 0.006, 0.0036
BS 125, conv only, reverse ramp	conv: 0.017, 0.01, 0.01, 0.01, 0.0036

Observations

Targets (held-out seeds 100-107): baselines 93.92 (BS 2000), 94.00 and 93.86 (BS 500 and 125, from the post-campaign reruns; the table's own baseline runs replace these once they finish). One held-out mean has a standard error of about 0.045 pp, so differences below about 0.1 pp are not resolved.

Conv only without a transform, v1 lands 0.1-0.2 pp short at every batch size: 93.82 (both search seeds) at BS 2000, 93.84 at BS 500 and 93.73 at BS 125, for 9.4-10 full runs of train steps. Its BS 2000 schedules (0.13, 0.13, 0.13, 0.047, 0.006 and 0.22, 0.13, 0.078, 0.047, 0.01) are one or two grid steps from C4 (93.95).

Under transforms the rest-of-run linear lookahead fails: the ramp gives 93.15 (BS 2000) and 93.40 (BS 500), below the do-nothing 93.42 / 93.37, and the wave 93.54 / 93.50. The lookahead decays linearly in the spec schedule, so under a transform its applied shape is distorted: under the ramp the applied LR rises late in the lookahead, so high early candidates look bad and the early intervals stay at 0.13-0.22 where the compensating spec is about 0.9. The reverse ramp, which mostly needs smaller early values, is fine (93.80 / 93.88).

From scratch with all four LRs (b), the level scan fails: 91.65-93.70 at BS 2000 and 92.93-93.58 at BS 500. Scored after one interval, head and norm_bias look flat (87.3% at levels 0.13, 1 and 7.7 in one search), so the scan leaves them near their start of 1.0, where a full run loses about 1 pp for head (Experiment 7). The search seed 1 run without a transform (92.01) kept head and norm_bias at 1.0.

The search's own final score is optimistic by up to 0.2 pp at BS 500 and 125 (e.g. 94.06 own vs 93.92 held-out; 93.94 vs 93.76): its decisions are selected on one seed.

v1 was stopped after these cells (the remaining transformed and 4-LR cells were cancelled) in favour of v2.

Experiment 14. ICS v2: a short-cooldown pass, then a rest-of-run pass on its plan

Question

v1's rest-of-run linear lookahead is defined on the spec schedule, so under a transform its applied shape is distorted and it misjudges early intervals (v1 under the ramp: 93.11-93.19, below doing nothing). Does a first pass with a short, local cooldown, whose schedule then serves as the second pass's lookahead, fix that?

Setup

Item	Value
Pass 1	Every interval in order: climb conv from the previous interval's value, each candidate scored after the interval plus one interval at 0.22 x its value (the least biased short cooldown, Experiment 12; local, so a transform's global shape does not distort it)
Pass 2	From the initial state again: each interval's climb starts at pass 1's value; a candidate is scored after the rest of the run on pass 1's schedule, its tail scaled by candidate / pass-1 value (rest_plan); the last interval by the final TTA
Other	As v1 (5 constant intervals, x0.6 grid, budget 20 full runs, 2 search seeds, held-out seeds 100-107); conv only: v2's 4-LR variant is replaced by v3/v4 (level rule)

Held-out TTA (mean over search seeds; vs the baseline; per search seed); cost in full runs

Batch size, scenario (baseline)	none	ramp	reverse ramp	wave
BS 2000, conv only (93.92)	**93.70** (-0.22; 93.77/93.62); 16.2x	**93.72** (-0.20; 93.68/93.77); 19.8x	**93.77** (-0.15; 93.83/93.71); 16.6x	**93.68** (-0.24; 93.84/93.53); 16.0x
BS 2000, 4 LRs (93.92)	running	running	running	running
BS 500, conv only (94.00)	**93.88** (-0.11; 93.89/93.88); 16.6x	**93.61** (-0.39; 93.54/93.68); 20.2x*	**93.81** (-0.19; 93.81/93.81); 16.6x	**93.78** (-0.22; 93.70/93.85); 15.6x
BS 500, 4 LRs (94.00)	running	running	running	running
BS 125, conv only (93.86)	**93.47** (-0.39; 93.43/93.52); 16.0x	**93.37** (-0.49; 93.50/93.25); 18.0x	**93.75** (-0.11; 93.72/93.78); 18.8x	**93.55** (-0.31; 93.40/93.70); 15.6x
BS 125, 4 LRs (93.86)	running	running	running	running

Cells within 0.1 pp of the baseline (about 2 SE of an 8-seed mean) or above: 0 of 12 finished. * = the search hit its step budget. Do-nothing: the baseline schedule kept under the transform.

Committed schedules (search seed 0)

Cell	Committed schedules (search seed 0)
BS 2000, conv only, none	conv: 0.13, 0.13, 0.078, 0.078, 0.017
BS 2000, conv only, ramp	conv: 0.36, 0.22, 0.047, 0.017, 0.0036
BS 2000, conv only, reverse ramp	conv: 0.078, 0.078, 0.078, 0.028, 0.017
BS 2000, conv only, wave	conv: 0.078, 0.078, 0.078, 0.047, 0.01
BS 500, conv only, none	conv: 0.078, 0.047, 0.047, 0.017, 0.0036
BS 500, conv only, ramp	conv: 0.22, 0.078, 0.017, 0.017, 0.0022
BS 500, conv only, reverse ramp	conv: 0.047, 0.047, 0.028, 0.028, 0.017
BS 500, conv only, wave	conv: 0.047, 0.028, 0.017, 0.028, 0.006
BS 125, conv only, none	conv: 0.028, 0.017, 0.006, 0.01, 0.0022
BS 125, conv only, ramp	conv: 0.13, 0.028, 0.006, 0.0036, 0.00047
BS 125, conv only, reverse ramp	conv: 0.028, 0.017, 0.028, 0.017, 0.006
BS 125, conv only, wave	conv: 0.017, 0.006, 0.01, 0.006, 0.0036

Observations

The short-cooldown first pass fixes v1's transform failure: under the ramp the held-out TTA rises from 93.15 (v1) to 93.72 at BS 2000 and from 93.40 to 93.61 at BS 500, because pass 1's local cooldown lets the early intervals rise (0.36-0.6 at BS 2000, where v1 stayed at 0.13-0.22) and pass 2 then judges candidates against a lookahead with the right, transform-shaped profile.

It is still 0.1-0.5 pp short everywhere: 93.68-93.77 at BS 2000 (baseline 93.92), 93.61-93.88 at BS 500 (94.00) and 93.37-93.75 at BS 125 (93.86), for 15-20 full runs. At BS 125 it is worse than v1 without a transform (93.47 vs 93.73): pass 1 decays faster than v1's rest-of-run lookahead (0.028, 0.017, 0.006, ... vs 0.028, 0.028, 0.017, ...) and pass 2, starting each climb from pass 1's values, does not undo it.

Search-seed spread is large (up to 0.31 pp between the two seeds of one cell), so single-seed decisions still dominate the error (Experiment 17).

Experiment 15. ICS v5: v2 plus kept first interval and replicated late decisions

Question

The late intervals' decisions are worth ~0.1 pp each and one seed gets them wrong about half the time (Experiment 12), while the first interval's rest-of-run candidates cost a full run each. Does keeping pass 1's first interval and scoring the late candidates by several continuations close the remaining gap within budget?

Setup

Item	Value
As v2	Pass 1 with the 0.22 x short cooldown, pass 2 with the rest_plan lookahead
Kept interval	Pass 2 keeps pass 1's first interval (value and trained state, at no cost)
Replicates	In pass 2, a candidate for interval 4 is scored by the mean of 2 continuations and for interval 5 of 3: the first on the run's own data stream (its state is committed), the others restore the pre-interval state, reseed the data order and augmentation (fresh epoch) and train interval + lookahead again
4 LRs	Levels: x7.7 scan right after interval 1 (conv: best score; head and biases: largest within 1 pp of the best), then as conv only with the other groups tied to conv's profile

Held-out TTA (mean over search seeds; vs the baseline; per search seed); cost in full runs

Batch size, scenario (baseline)	none	ramp	reverse ramp	wave
BS 2000, conv only (93.92)	**93.75** (-0.17; 93.95/93.56); 14.6x	**93.58** (-0.34; 93.66/93.51); 17.0x	**93.81** (-0.11; 93.91/93.71); 18.6x	**93.79** (-0.13; 93.84/93.75); 15.4x
BS 2000, 4 LRs (93.92)	running	running	running	running
BS 500, conv only (94.00)	**93.86** (-0.14; 93.83/93.89); 16.2x	**93.65** (-0.35; 93.62/93.68); 19.6x*	**93.81** (-0.19; 93.81/93.81); 16.0x	**93.76** (-0.24; 93.68/93.84); 15.0x
BS 500, 4 LRs (94.00)	running	running	running	running
BS 125, conv only (93.86)	running	running	running	running
BS 125, 4 LRs (93.86)	running	running	running	running

Cells within 0.1 pp of the baseline (about 2 SE of an 8-seed mean) or above: 0 of 8 finished. * = the search hit its step budget. Do-nothing: the baseline schedule kept under the transform.

Committed schedules (search seed 0)

Cell	Committed schedules (search seed 0)
BS 2000, conv only, none	conv: 0.13, 0.13, 0.078, 0.047, 0.01
BS 2000, conv only, ramp	conv: 0.36, 0.22, 0.047, 0.01, 0.0022
BS 2000, conv only, reverse ramp	conv: 0.078, 0.078, 0.078, 0.078, 0.028
BS 2000, conv only, wave	conv: 0.078, 0.078, 0.078, 0.047, 0.006
BS 500, conv only, none	conv: 0.078, 0.047, 0.047, 0.028, 0.0036
BS 500, conv only, ramp	conv: 0.22, 0.078, 0.017, 0.01, 0.0013
BS 500, conv only, reverse ramp	conv: 0.047, 0.047, 0.028, 0.028, 0.017
BS 500, conv only, wave	conv: 0.047, 0.028, 0.017, 0.017, 0.006
BS 125, conv only, none	conv: 0.028, 0.017, 0.006, 0.0036, 0.0013
BS 125, conv only, ramp	conv: 0.078, 0.017, 0.01, 0.0036, 0.00047
BS 125, conv only, reverse ramp	conv: 0.017, 0.01, 0.01, 0.006, 0.006

Experiment 16. How well do constant intervals represent the baseline?

Question

Before asking a search to find a constant-interval schedule, can such a schedule match the baseline at all, at each batch size and under each transform, with 5 or 10 intervals?

Setup

Item	Value
Schedules	conv's schedule replaced by n equal constant intervals, each at the oracle spec (the baseline's applied LR divided by the transform) a fraction q of the way through the interval; the last interval at min(q, 0.95) (the oracle reaches 0 at the end). q = 0.5 is about the interval mean, q = 1 the end value. Exact values (no grid); the other groups on the baseline
Runs	Plain training runs on held-out seeds 100-107 (one standard error about 0.045 pp)

Held-out TTA of the oracle staircases

Batch size, transform	n = 5: q 0.5 / 0.75 / 1	n = 10: q 0.5 / 0.75 / 1	Baseline (do-nothing)
BS 2000, none	- / - / -	- / - / -	93.92
BS 2000, ramp	- / - / -	- / - / -	93.92
BS 2000, reverse ramp	- / - / -	- / - / -	93.92
BS 2000, wave	- / - / -	- / - / -	93.92
BS 500, none	- / 93.85 / 93.70	- / 94.00 / 93.88	94.00
BS 500, ramp	- / 93.75 / 93.68	- / 94.05 / 93.89	94.00
BS 500, reverse ramp	- / 93.91 / 93.62	- / 93.88 / 93.84	94.00
BS 500, wave	- / 93.86 / 93.59	- / 93.89 / 93.85	94.00
BS 125, none	- / 93.80 / 93.72	- / 93.81 / 93.84	93.86
BS 125, ramp	- / 93.72 / 93.57	- / 93.90 / 93.84	93.86
BS 125, reverse ramp	- / 93.84 / 93.57	- / 93.86 / 93.71	93.86
BS 125, wave	- / 93.78 / 93.55	- / 93.95 / 93.82	93.86

Experiment 17. Averaging two searches' schedules

Question

Two search seeds of the same ICS version often commit different schedules, one or two grid steps apart in a few intervals, and score 0.1-0.4 pp apart. Their decision errors are partly independent; does the geometric mean of the two schedules, interval by interval, score better than the two do on average?

Setup

Item	Value
Schedules	For each finished cell of ICS v1, v2 and v5 (conv only): the geometric mean of search seeds 0 and 1's committed conv values, interval by interval (rounded to 2 significant figures, not snapped to the grid)
Runs	Plain training runs on held-out seeds 100-107

Held-out TTA

Cell	Averaged	Parents (seed 0 / 1)	Gain over their mean	Baseline
v1, BS 125, none	**93.77**	93.70 / 93.76	+0.04	93.86
v1, BS 125, reverse ramp	**93.76**	93.72 / 93.60	+0.10	93.86
v1, BS 2000, none	**93.90**	93.82 / 93.82	+0.08	93.92
v1, BS 2000, ramp	**93.16**	93.19 / 93.11	+0.02	93.92
v1, BS 2000, reverse ramp	**93.89**	93.81 / 93.79	+0.10	93.92
v1, BS 2000, wave	**93.63**	93.61 / 93.46	+0.09	93.92
v1, BS 500, none	**93.81**	93.75 / 93.92	-0.03	94.00
v1, BS 500, ramp	**93.43**	93.27 / 93.53	+0.04	94.00
v1, BS 500, reverse ramp	**93.95**	93.95 / 93.82	+0.07	94.00
v1, BS 500, wave	**93.66**	93.25 / 93.75	+0.16	94.00
v2, BS 125, none	**93.66**	93.43 / 93.52	+0.18	93.86
v2, BS 125, ramp	**93.49**	93.50 / 93.25	+0.12	93.86
v2, BS 125, reverse ramp	**93.74**	93.72 / 93.78	-0.01	93.86
v2, BS 125, wave	**93.70**	93.40 / 93.70	+0.15	93.86
v2, BS 2000, none	**93.81**	93.77 / 93.62	+0.12	93.92
v2, BS 2000, ramp	**93.74**	93.68 / 93.77	+0.02	93.92
v2, BS 2000, reverse ramp	**93.78**	93.83 / 93.71	+0.01	93.92
v2, BS 2000, wave	**93.78**	93.84 / 93.53	+0.10	93.92
v2, BS 500, none	**93.90**	93.89 / 93.88	+0.02	94.00
v2, BS 500, ramp	**93.66**	93.54 / 93.68	+0.05	94.00
v2, BS 500, reverse ramp	**93.92**	93.81 / 93.81	+0.11	94.00
v2, BS 500, wave	**93.82**	93.70 / 93.85	+0.04	94.00
v5, BS 2000, none	**93.75**	93.95 / 93.56	-0.01	93.92
v5, BS 2000, ramp	**93.62**	93.66 / 93.51	+0.04	93.92
v5, BS 2000, reverse ramp	**93.86**	93.91 / 93.71	+0.06	93.92
v5, BS 2000, wave	**93.85**	93.84 / 93.75	+0.06	93.92

Mean gain over the parents' mean: +0.066 pp over 26 cells (23 positive).

Observations

Averaging two search seeds' schedules helps on average: the averaged schedule beats the mean of its two parents by +0.066 pp over 26 cells (v1, v2 and v5, conv only), positive in 23 of them, by up to +0.18 pp (v2 at BS 125). It mostly recovers the better parent: when one seed made a costly mistake, the average sits between the two choices, and the score curves are flat enough near the optimum that the midpoint costs little.

This is decision noise being averaged out: the search's single-seed decisions err by a grid step in some intervals (Experiment 12), and two independent searches err in different intervals. Two v1 searches averaged (about 20 full runs together) reach 93.90 at BS 2000 without a transform and 93.89 under the reverse ramp, against 93.75 and 93.81 for one v5 search (about 15 full runs); averaging cannot fix a shared bias, though (v1 under the ramp stays at 93.16).