


Input-conditioned head grid | 120/120 trials
Conv: muon2 | batch size 2000 | 8 epochs | momentum after conditioning

Linear decay | momentum version 3 | Nesterov True
Best 94.02% | LR 10000, momentum 0.8



Linear decay | momentum version 3 | Nesterov False
Best 94.10% | LR 10000, momentum 0.7

